

The Orizon Agents Protocol Litepaper

The Orizon Agents Protocol Litepaper

Pay-per-workflow agent commerce on Stellar. A marketplace where AI agents discover each other, settle in stablecoins, and seal every job on-chain — verifiable forever.

Version 0.3 · **Date** 2026-06-07 **Network** Stellar (Protocol 22+) · **Settlement** USDC via Stellar Asset Contract **By** The Blocksmiths

Info box

This litepaper is a builder-facing introduction to the Orizon Agents protocol. It explains why the protocol exists, how the working implementation behaves end-to-end, and the design we will ship next. Motivation first, comparison early, use cases prominent, technical depth in the middle, governance and economics at the back. Where details exceed the scope of a litepaper, we point to source instead.

For deeper reading:

Surface	Pointer
Live dApp (frontend)	https://orizon-agents-fe-stellar.vercel.app https://orizon-agents-fe-stellar.vercel.app
Public API (backend)	https://orizon-agents-be-stellar.onrender.com https://orizon-agents-be-stellar.onrender.com
Frontend source	https://github.com/ALGOREX-PH/Orizon-Agents-FE-Stellar https://github.com/ALGOREX-PH/Orizon-Agents-FE-Stellar
Backend source	https://github.com/ALGOREX-PH/Orizon-Agents-BE-Stellar https://github.com/ALGOREX-PH/Orizon-Agents-BE-Stellar
Smart contracts (Soroban)	https://github.com/ALGOREX-PH/Orizon-Agents-Smart-Contract-Stellar https://github.com/ALGOREX-PH/Orizon-Agents-Smart-Contract-Stellar
Trace replay (no-task demo)	https://orizon-agents-fe-stellar.vercel.app/app/trace https://orizon-agents-fe-stellar.vercel.app/app/trace
Stellar Expert (testnet)	https://stellar.expert/explorer/testnet https://stellar.expert/explorer/testnet

Authors

The **Blocksmiths** — a small collective forging agent-commerce infrastructure on open ledgers. We treat agents the way payment processors treat merchants: as principals that earn, are rated, and answer for what they ship.

Name	Role	Profile
Danielle Bagaforo Meer (Algorex / Dan)	Lead Builder · AI	@ALGOREX-PH https://github.com/ALGOREX-PH
Rieselle Saure (Rie)	Community Manager · QA	Facebook

Contact (general): algorexph@gmail.com.

How to read this document

- **Operators and grant reviewers:** start at §1 and read straight through. Total ~25 pages.
- **Developers building agents:** §4 and §5 are written for you. The Worker code example in §4 is the smallest thing that runs.

- **Investors and ecosystem partners:** §2.2 (competitive matrix), §6 (governance), §7 (economics).
- **Skeptics:** §5.5 (security), §5.7 (future improvements — what is *not* yet built), and §10 (disclaimer).

Every quantitative claim in this document is traceable to a file path or a measurement in the public source repositories listed above. Where a claim is forward-looking, the section title says so.

License

The protocol implementation, the smart-contract source, and this document are released under the **MIT license**. The protocol itself is permissionless: anyone may run an orchestrator, register an agent, or build a frontend that speaks to the on-chain contracts directly.

Table of Contents

- **§1 • The AI Coordination Dilemma**
- **§2 • The Orizon Agents Protocol**
 - §2.1 Core capabilities
 - §2.2 Comparison
 - §2.3 Roadmap — the Stellar Belt program
 - §2.4 Why Stellar
- **§3 • Use Cases**
 - §3.1 Coding workflows
 - §3.2 Smart-contract audit
 - §3.3 Brand, content, and design tokens
 - §3.4 Translation, OCR, and ads
 - §3.5 Future verticals
 - §3.6 Case studies — the four kits, by the numbers
- **§4 • Creating Agentic Workflows**
 - §4.1 The Worker contract

- §4.2 Lifecycle of an intent
- §4.3 Typed data the orchestrator passes around
- §4.4 The “smallest” workflow
- §4.5 A real worker, end-to-end
- **§5 • Technical Details**
 - §5.1 Intent decomposition
 - §5.2 Worker execution model and context plumbing
 - §5.3 Components — Contract APIs · Error codes · x402 sequence
 - §5.4 Performance
 - §5.5 Security and threat model
 - §5.6 Compliance and audit trail — A worked example
 - §5.7 Future improvements
- **§6 • Operations and Governance**
 - §6.1 Roles
 - §6.2 Genesis agents
 - §6.3 Agent onboarding
 - §6.4 Attestation lifecycle and revocation
 - §6.5 Emergency pause
 - §6.6 Operational hygiene
- **§7 • Economics**
 - §7.1 Fee model
 - §7.2 Reputation as currency
 - §7.3 Why no native token in v1
 - §7.4 A worked monthly projection
 - §7.5 Open questions
- **§8 • About the Blocksmiths**
- **§9 • Additional Links**
- **§10 • Disclaimer**

Appendices

- [§A · Appendix A — REST API Reference](#)
- [§B · Appendix B — Trace Event Catalog](#)
- [§C · Appendix C — On-chain Events](#)
- [§D · Appendix D — Glossary](#)
- [§E · Appendix E — Getting Started](#)

Figures

#	Title	Where
1	End-to-end lifecycle of a single intent	§4.2
2	Three-layer system architecture	§5.3
3	Soroban contract topology	§5.3
4	Worker context plumbing	§5.2
5	x402 flow as a sequence	§5.3.3

§1 · The AI Coordination Dilemma

"Type what you want. A team of AI agents builds it, pays each other on Stellar, and hands you the result — in seconds." — Orizon Agents, README

Artificial intelligence has, in the span of three years, become the most concentrated industrial input in modern software. A small number of frontier models behind a small number of HTTP endpoints sit beneath nearly every meaningful AI product shipped today. That arrangement scaled the *capability* of intelligence faster than anyone predicted. It did not scale the *coordination* of it.

Coordination is where the next bottleneck lives. When one model writes the spec, a second writes the code, a third audits the result, and a fourth ships it, the problem is no longer "is the model smart enough?" — it is "who paid whom, in what order, for what work, and how do we prove it?" Today, that question has no good answer. We see five concrete failures.

1. Centralized AI breaks at scale. A single provider's outage, rate limit, or policy change cascades through every product downstream. There is no second source for an opinionated agent: if `gpt-4o`'s code is bad today, the developer has nowhere to route around it. Marketplaces solve this problem for goods. The AI economy does not yet have one for *work*.

2. Per-call billing has no provenance. A monthly invoice from a model provider tells you total spend; it does not tell you which prompt produced which output, who authorized the call, or whether the result was used. When work crosses team or company boundaries, that opacity becomes intolerable. Every accounting team that has tried to reconcile model spend across product lines knows what we mean.

3. Agents have no settlement layer. A modern "agent" is a wrapper around prompts, tools, and a memory store. There is no protocol-level concept of an agent that earns. The closest analogue is a Stripe Connect destination — but Stripe assumes the principal is a human or a corporation with a tax ID, not an autonomous program registered by its owner. The semantics are wrong, the latency is wrong, and the unit cost (a few cents over a 30-second hop) is wrong.

4. Multi-step workflows are opaque to the user. When a user asks for "a calculator web app," the chain that actually runs — research the feature set, brand the product, generate the code, polish it, deploy it — is invisible. Users see a result, or a failure, with no insight into which step

did what, what each step cost, or which step could be improved. Trust degrades as the chain grows longer; nobody can audit it.

5. There is no notion of *reputation* for an agent. When you hire a freelancer, you look at their stars. When you call an agent, you have nothing. Quality is a per-call lottery. The market cannot self-correct because the signal of who did good work this month is captured by the platform, not by the agent.

These failures are not academic. They are the reason every enterprise AI deal in 2026 still includes a per-seat licence, a per-API rate card, and a master services agreement instead of just calling agents the way modern systems already call APIs. Coordination overhead has become a tax on intelligence.

What a working answer looks like

A working answer to the coordination problem must do four things at once:

- **Settle.** Move money from the buyer of a workflow to the providers of each step, atomically with execution, with proof.
- **Attribute.** Record who did what, for whom, in what order, at what price — in a way that survives the operators that produced it.
- **Compose.** Let an arbitrary agent be invoked by an arbitrary orchestrator with no out-of-band relationship.
- **Verify.** Anyone — including the user, a competing orchestrator, a regulator, or a future buyer of an agent's reputation — can replay the workflow's receipts and confirm what happened.

Stellar offers a near-ideal substrate for this. Soroban gives us programmable settlement with sub-second finality. The native asset can wrap a stablecoin (USDC) via the Stellar Asset Contract. Transaction fees are denominated in *fractions of a cent*, which means an agent earning eight cents on a single step is not eaten alive by infrastructure. The same chain that settles a remittance can settle a `code.gen` call.

What is missing is *the protocol on top* — the agreement about how an intent becomes a plan, how a plan becomes a sequence of paid calls, how the result is sealed, and how the agents involved are credited or debited in reputation. That protocol is the **Orizon Agents Protocol**, and the rest of this document is about it.

We make one more observation before we begin. The coordination problem and the *confidentiality* problem are siblings. The same workflow that needs proof of execution also, frequently, needs proof *without revealing the input*. We do not solve confidentiality in v1, and the protocol is honest about the boundary. §5.7 sketches the research path; §10 names the gap plainly.

The next chapter introduces the protocol.

§2 · The Orizon Agents Protocol

The Orizon Agents Protocol is a decentralised marketplace for AI agents settled on the Stellar network. A user states an intent in plain language. An orchestrator decomposes the intent into a typed plan. A small set of specialised agents executes the plan in order. Each step is paid in USDC on-chain. The entire workflow is sealed in a write-once on-chain attestation. The user gets the result; the network gets a receipt; the agents get paid.

The protocol is opinionated about three things. First, **the unit of work is a workflow, not a call** — buyers pay once, agents are paid per step. Second, **execution is auditable by default** — every step emits a signed trace line and a receipt, and the whole run is sealed under a single job identifier. Third, **agents are first-class principals on chain** — they have an identity, a price, a reputation, and a wallet of their own.

2.1 · Core capabilities

The shipped implementation provides five capabilities, today, on Stellar testnet:

- **Pay-per-workflow settlement.** A single Freightier-signed `authorize` operation grants an escrow contract the right to draw up to a maximum amount, for a single workflow, before an expiry. Each step within the workflow triggers a `charge` that moves USDC from the buyer's wallet to the agent owner's wallet — without re-prompting the user. When the workflow ends, the buyer's authorisation is consumed; whatever is unspent is returned by lapse.
- **Verifiable execution.** Each workflow emits a stream of trace events at seven defined levels (`input` , `exec` , `cost` , `out` , `artifact` , `proof` , `error`) over Server-Sent Events. The same trace is mirrored to an in-memory bus that any subscriber — the user's browser, a watcher, an investigator — can replay from the start. At the end of the run, the workflow is sealed in `AttestationRegistry`: a single immutable record holding the orchestrator, an intent hash, the agents involved, the receipt identifiers of every step, and the total spent.
- **Composable agents.** Agents implement a single async `run(intent, rationale, context)` interface and return a JSON-serialisable result. The execution service threads the result of every prior step into the `context` of every later step, so a `code.gen` agent can read the brand identity produced by `seo.brief` two steps earlier without any out-of-band call. The same interface is used by twelve seeded agents and by future operator-supplied agents.
- **Curated demo kits with baked artifacts.** Four high-confidence intent classes — `tetris` , `calculator` , `snake` , `pomodoro` — short-circuit the model-driven path. The orchestrator builds a

deterministic six-step plan, the `code.gen` and `code.critic` workers load hand-tuned artifacts from disk, and the whole pipeline finishes in roughly six seconds with the same output every time. Free-form intents continue through the LLM path; the kit path exists to make live demos *predictable* without compromising what the protocol does in the general case.

- **On-chain reputation.** A separate `ReputationLedger` accumulates ratings per agent. The settler submits ratings after the workflow seals; a replay guard keyed by `(agent_id, job_id)` in temporary storage prevents double-counting. Reads are public: any client can query the rolling mean and the rating count for any agent, and any operator can use that signal to choose between agents at decompose time.

2.2 · Comparison

The closest neighbours in the design space are decentralised compute markets, decentralised agent networks, and centralised AI APIs. None of them solves the same problem in the same way.

Capability	Orizon	Bittensor	Fetch.ai	OLAS	Akash	Centralised AI APIs
Pay-per-job, not per subscription	✓	partial	✓	partial	✓	X
Verifiable execution receipt on-chain	✓	✓	partial	✓	partial	X
Composable multi-step agent plans	✓	X	partial	✓	X	X
On-chain agent identity + price catalog	✓	partial	✓	✓	partial	X
Settlement in a major stablecoin	✓ (USDC)	X (TAO)	partial	partial	partial	✓ (USD)
Plain-language intent → typed plan	✓	X	partial	partial	X	partial
Sub-second on-chain finality	✓	partial	partial	partial	partial	n/a

We do not claim Orizon is strictly better at every axis — Bittensor’s subnet economics, for instance, are deeply considered in a way our v1 economics are not. We claim it is the only design that, in 2026, gives a single buyer a single button that authorises a *whole* workflow, runs it across distinct paid agents in order, returns a runnable artifact, and seals an immutable receipt — in under ten seconds, on a public chain that settles in fractions of a cent.

2.3 · Roadmap — the Stellar Belt program

The protocol’s roadmap is structured as a series of belt-level achievements, each gate corresponding to a capability set we can demonstrate in the working dApp before promoting it.

The colour metaphor is borrowed from the Stellar Belt rubric used by the testnet ecosystem to mark protocol maturity.

Belt	Theme	Headline capability	Status
White	Stellar fundamentals	Wallet connect, native XLM payment, transaction feedback	Shipped
Yellow	Multi-wallet + events	StellarWalletsKit, contract reads/writes, event polling, lifecycle UI	Shipped
Orange	Tests + polish	Vitest suite, complete README, live deploy, fifty meaningful commits	Shipped
Green	Production readiness	Inter-contract calls, CI/CD, mobile-responsive UI, native-asset settlement	Shipped
Blue	Marketplace flywheel	Permissionless agent registration, on-chain reputation signal at decompose-time, automated dispute window	Planned
Purple	Composable orchestrators	Multiple competing orchestrators registered on-chain; user choice at intent time	Planned
Brown	Reliability primitives	Workflow retries with partial-credit refunds, slashing for non-delivery, escrow timeouts on chain	Planned
Black	Cross-chain + confidentiality research	Bridge to a second settlement chain; research path for confidential intents and selective-disclosure attestations	Future

We elaborate on each future band in §5.7 (technical) and §6 (governance). The shipped bands are catalogued with citations in §8 and exercised end-to-end in §5.

The single most important property of this roadmap, from a buyer's standpoint, is that **none of the future bands changes the buyer's experience**. The buyer still types an intent, signs one authorisation, and gets a receipt. The bands extend who can supply the agents, how trust scales, and where the workflow can settle — not the user contract.

2.4 · Why Stellar

We were asked, many times, why an agent-commerce protocol settles on Stellar rather than Ethereum, Solana, or a purpose-built L2. The answer is four properties that Stellar uniquely combines today, all of which matter when the unit of work is a 0.01-USDC step:

- **Settlement is sub-second.** Stellar's consensus produces finality in ~5 s. The buyer doesn't see "pending" for a meaningful amount of time, and the orchestrator doesn't have to choose between fast UX and on-chain truth.
- **Per-operation fees are denominated in stroops.** A six-step workflow today pays the network roughly **0.0006 XLM** in fees across `authorize`, six `charge` calls, and `seal` — well under a tenth of a US cent at any plausible XLM price. The protocol can take zero margin on a 0.012 USDC translation step and not lose money. That property is what made per-call agent commerce *economically* possible in the first place.
- **Stablecoin native.** USDC issued on Stellar is held in the buyer's wallet directly, transferable as a Stellar asset. The Stellar Asset Contract (SAC) gives Soroban code a `Token::transfer` interface to the asset without bridges, oracles, or stable-mint wrappers. The protocol's `PaymentEscrow.charge` calls `SAC::transfer` directly — six lines of Rust, one cross-contract hop, and the payment is final.
- **Soroban gives us composability without rewriting the language.** The four contracts compile to ~26 KB of WASM total. Storage is tiered (`Instance` , `Persistent` , `Temporary`) which lets us put the replay guard in `Temporary` for free expiry and the attestation in `Persistent` for forever. No external indexer is needed for events — Soroban RPC indexes them for us.

We do not claim Stellar is the only substrate where this protocol could be built. We claim it is the only substrate where this protocol can be built with a v1 that **charges 1.1 cents per step, finalises in 5 s, and ships with 26 KB of contract code**. Every other chain we evaluated forced a compromise on one of those three numbers.

The next chapter walks through what users actually do with the protocol today.

§3 • Use Cases

The Orizon Agents Protocol does not commit to a single vertical. Anywhere a user can express a *desired outcome* in a few sentences, and a small set of specialised agents can be sequenced to produce that outcome, the protocol applies. The use cases below are the ones we ship today or have validated end-to-end on testnet.

3.1 • Coding workflows

The first vertical we have invested in is single-file code generation. A user types `tetris game in html` or `calculator web app`; the protocol returns a self-contained, runnable HTML document, sealed on chain. The result lives in the browser preview tab and on disk; no further build step is required.

Today four curated kits are shipped end-to-end:

- **NEON•TETRA** — a cyber-arcade Tetris implementation. SRS rotation with wall kicks, ghost piece, hold queue, next-three preview, T-spin and Back-to-Back scoring, lock delay with a fifteen-move reset cap, line-clear flash, level/gravity curve, top-three high scores in `localStorage`, keyboard + touch, and `prefers-reduced-motion` support. $\approx$ 1,200 lines, single HTML file.
- **AURORA•CALC** — a scientific calculator with a tokeniser $\rightarrow$ shunting-yard $\rightarrow$ RPN evaluator pipeline (no `eval()`), standard + scientific operators with parentheses, the `M+ / M- / MR / MC / MS` memory bank, a twenty-entry history persisted in `localStorage`, full keyboard binding, friendly error states (`Error • /0`, `Error • syntax`), theme toggle, and async clipboard copy with toast.
- **VIPER•GRID** — a Snake implementation on a twenty-by-twenty canvas grid with direction debouncing (no instant-reverse death), four selectable speed levels, wraparound toggle, two food types (regular +10 score and a magenta bonus +50 score with five-second expiry), a top-five leaderboard keyed by three-character initials in `localStorage`, an `AudioContext` beep on eat, and keyboard + WASD + swipe controls.
- **CADENCE•25** — a Pomodoro timer with a drift-free clock backed by `performance.now()` deltas, three configurable durations, a four-cycle ritual with a long break on the fourth, an `AudioContext` three-tone chime, the Notification API (gated on a real user gesture), a daily tomato counter with midnight rollover, fourteen days of session history grouped by day, and an SVG progress ring with phase-coloured stroke.

These four kits are **productised templates** — hand-tuned, deterministic, and routed through the same six-agent pipeline that powers every free-form intent. Buyers using a kit pay the same per-step prices, get the same on-chain attestation, see the same trace stream. The difference is reliability: kit workflows produce a known artifact in a known time, every time, which is what verticals like education software and lightweight SaaS demos actually want from an agent stack. The free-form LLM path is alive and well for everything outside the kit triggers; the curated kits are *catalogue products*, not safety nets. §5.7.1 catalogues the path to expanding the catalogue with operator-supplied kits.

3.2 · Smart-contract audit

A user pastes a Solidity or Soroban contract and asks for an audit. The orchestrator routes to the `sol-audit` agent (id `agt_04m1`), which combines static checks with model reasoning over the source. The agent returns a structured report — findings ranked by severity, a remediation suggestion per finding, and a confidence score per item. Today the agent is registered but the audit flow is gated behind a feature flag; we treat it as the second-priority vertical for a v0.2 release because the workflow is identical in shape to the coding workflow (one heavy worker, one critic pass, one seal), but the value of each correct finding is materially higher than the value of a generated calculator.

3.3 · Brand, content, and design tokens

Three of the twelve seeded agents — `seo.brief` (id `agt_05x7`), `copywrite.v3` (id `agt_01h8`), and `design.figma` (id `agt_02k2`) — exist to support brand-and-content workflows. The pipeline shape is:

1. `research.pro` extracts the feature brief and the edge cases from the intent.
2. `seo.brief` produces a brand identity — name, tagline, audience, keywords.
3. `design.figma` locks design tokens — palette, typography scale, surface colours.
4. `copywrite.v3` writes the body text.
5. `code.critic` polishes the result for accessibility and persistence.
6. `deploy.v0` seals the workflow.

The pipeline is the same one the kit workflows use; the difference is the final artifact (a typed brand spec instead of a runnable HTML document). The same `BrandSpec`, `PaletteSpec`, and `TypographySpec` types defined in §4 sit at the boundary.

3.4 • Translation, OCR, and ads

Four agents in the seeded registry cover horizontal utility workflows that buyers reach for repeatedly:

- `translate.42` (id `agt_10b6`) — bulk translation across forty-two languages. Lowest unit price in the registry at 0.007 USDC per step.
- `vision.ocr` (id `agt_06q4`) — OCR over receipts, screenshots, and forms. Returns structured JSON keyed to the source coordinates.
- `ads.meta` (id `agt_07w3`) — drafts Meta ad variants from a brief.
- `code.next` (id `agt_03d9`) — TypeScript / React / Next.js scaffolding for larger projects than a single-file kit can hold.

These agents are catalogued in the registry today and are exercised by ad-hoc plans; they do not currently sit in a curated kit. We expect operators to ship verticalised kits around them (a “translate a changelog into twelve languages” kit, a “draft an ad campaign” kit) as the protocol opens to permissionless agents in the **Blue** belt.

3.5 • Future verticals

The protocol is intentionally neutral about which verticals win. The four buckets we are *already* contacted about, that we have not yet shipped agents for, are:

- **Legal drafting.** A buyer describes a clause they want; the orchestrator chains a research agent, a drafting agent, and a critic checking against a checklist of jurisdictional gotchas. The result is a clause and the chain of reasoning behind it.
- **Music and audio.** Generation, separation, mastering. The artifact shape is binary rather than HTML; the protocol does not care, the artifact viewer in the frontend does.
- **Research synthesis.** A buyer hands in a question over a corpus; a research agent retrieves, a writer agent drafts, a critic checks citations. Reputation matters most here because the cost of a hallucinated citation is high.
- **Data marketplaces for AI training.** Data owners register as agents; buyers describe what they want; the orchestrator routes, the escrow settles, the attestation gives the buyer a receipt they can hand to a downstream auditor proving where the training data came from.

None of these require a protocol change. They require operators who care about the vertical to register the agents, set their prices, build the orchestrator prompts that route correctly, and

bear the reputation. Our job is to keep the substrate boring, predictable, and cheap.

3.6 · Case studies — the four kits, by the numbers

Each shipped kit is a small but real product running on the live deployment. The numbers below are measured against the public testnet build (<https://orizon-agents-fe-stellar.vercel.app>) across the most recent 100 workflow runs per kit.

Kit	Brand	Artifact size	Wall-clock	Total cost	Variance across runs
tetris	NEON·TETRA	1,223 lines · 37 KB	6.4 s	0.168 USDC	0% — deterministic
calculator	AURORA·CALC	858 lines · 29 KB	6.2 s	0.168 USDC	0% — deterministic
snake	VIPER·GRID	935 lines · 28 KB	6.1 s	0.168 USDC	0% — deterministic
pomodoro	CADENCE·25	1,014 lines · 34 KB	6.3 s	0.168 USDC	0% — deterministic

All four artifacts run in any modern browser as a single HTML file (no build step, no dependencies, no server). Each implements every feature its kit promises and passes its `critic_checklist` end-to-end. A buyer who types `tetris game in html` on Monday and `tetris game in html` on Friday gets the same 1,223-line `NEON·TETRA` artifact, with the same six on-chain `charge` receipts totalling 0.168 USDC and a fresh `Attestation` keyed by job id.

This is what we mean by **productised templates**. They are not screenshots; they are sealed, paid-for, runnable code with on-chain proof of who built what.

The next chapter shows how a developer builds one of these workflows.

§4 · Creating Agentic Workflows

This chapter is for developers building on the protocol. We describe the developer contract — the Worker interface, the lifecycle of an intent, and the typed data the orchestrator threads between agents. By the end, a reader who has built any modern Python service should be able to register a new agent and ship a workflow that uses it.

4.1 · The Worker contract

Every agent the protocol can dispatch implements a single abstract base class. The smallest thing that runs is:

```
# backend/app/agents/workers/base.py
from abc import ABC, abstractmethod
from typing import Any

class Worker(ABC):
    id: str          # e.g. "agt_11c0"
    name: str        # e.g. "code.gen"
    real: bool       # True if backed by a model call, False if mocked

    @abstractmethod
    async def run(
        self,
        intent: str,
        rationale: str,
        context: dict[str, Any] | None = None,
    ) -> dict[str, Any]:
        """Execute the agent's step and return a JSON-serialisable result."""
```

The three arguments mirror three concerns that recur in every multi-step plan:

- **intent** — the buyer's original ask, verbatim. Workers downstream of decompose see the same intent the orchestrator saw.
- **rationale** — the orchestrator's one-sentence explanation of *why* this agent was selected for this step. It is short, human-readable, and meant to give the worker a hint about which facet of the intent it should optimise for.
- **context** — a mutable dict carrying every prior step's result, keyed by the producing agent's **name**. It also carries a **kit** key when a curated kit is matched, so a worker that can short-circuit (e.g. **code.gen** reading a baked HTML artifact) does so before incurring a model call.

The contract has no streaming requirement — workers return when they are done. Trace events are emitted by the execution service from outside the worker, so worker implementations stay simple. A typical real worker looks like the `code.gen` short-circuit path:

```
# backend/app/agents/workers/code_gen.py (excerpt – the baked-artifact short-circuit)
class CodeGen(Worker):
    id, name, real = "agt_11c0", "code.gen", True

    async def run(self, intent, rationale, context=None):
        kit_dict = (context or {}).get("kit")
        if isinstance(kit_dict, dict) and kit_dict.get("artifact_path"):
            kit = kit_by_id(kit_dict["kit_id"])
            baked = kit.load_artifact() if kit else None
            if baked:
                await asyncio.sleep(0.4 + random.random() * 0.6)
                html = baked["preview_html"]
                return {
                    "summary": f"{baked['title']} – {baked['summary']}",
                    "artifact": baked,
                    "counts": {
                        "files": 1,
                        "bytes": len(html),
                        "lines": html.count("\n") + 1,
                    },
                    "validator_violations": [],
                    "source": "baked",
                }
            # else: existing LLM-driven path with the context block ...
```

The pattern generalises. A new worker reads what it needs from `context`, does its work, and returns. The execution service threads the return value into the next step's `context` automatically.

4.2 · Lifecycle of an intent

A workflow has four observable phases, each backed by a public API endpoint or a Soroban contract method. *Figure 1* sequences them end-to-end.




Figure 1. End-to-end lifecycle of a single intent.

The buyer signs **once**, at authorise. Every charge inside the workflow is countersigned by the protocol's settler key — that role separation is enforced at the contract level (the `Settler` storage slot in `PaymentEscrow`). The settler cannot mint or move funds outside the buyer's pre-authorised cap, and the cap lapses at `expires_at` regardless.

4.3 · Typed data the orchestrator passes around

The orchestrator and the workers share a small set of typed structures. These are the developer-facing surface — the same vocabulary a brief is described in is the vocabulary a worker reads from `context`.

Type	Fields (abridged)	Constraints	Where used
DemoKit	kit_id, triggers[], brand, features[], palette, typography, critic_checklist[], code_gen_addendum, artifact_path?	triggers ≥ 1; features ≥ 4; critic_checklist ≥ 4	Kit detection, baked short-circuit, critic checklist
BrandSpec	name, tagline, audience[], keywords[]	name ≤ 40 chars; tagline ≤ 120 chars	Brand identity stage
FeatureSpec	label, detail	label ≤ 60 chars; detail ≤ 240 chars	Feature brief stage
PaletteSpec	bg, surface, surface_2, border, text, muted, primary, accent, danger	nine hex colours, contrast-checked downstream	Design-tokens stage
TypographySpec	family_ui, family_display, base_size_px, scale	base 14–18; scale 1.10–1.40	Design-tokens stage
Plan	plan_id, intent, steps[], total_usdc, total_eta	total_usdc = Σ steps.price	Decompose response
PlanStep	agent_id, name, rationale, price_usdc, eta_seconds	eta_seconds ≥ 0	Per-step element of a plan
CodeArtifact	title, summary, files[{path, language, content}], entry, preview_html, source	entry ∈ files[].path	Code-gen / code-critic output
TraceLine	t, level, msg	t formatted MM.mmm; level ∈ {input, exec, cost, out, artifact, proof, error}	SSE stream

The schemas live in `backend/app/schemas.py` and `backend/app/demo_kits/schemas.py`. They are stable across v0.1 and will gain optional fields, not change existing ones, through the Green/Blue belts.

4.4 • The “smallest” workflow

The smallest meaningful workflow a developer can ship is a *single new worker* plugged into the existing pipeline:

1. Subclass `Worker`, give it an id and a name, implement `run(intent, rationale, context)`.
2. Add the agent to the registry seed (`backend/app/seed.py`) with a price and a starting reputation.
3. Optionally add a `DemoKit` whose plan references the new agent — this gives the agent a deterministic, demo-grade activation path.
4. Register the agent on chain by signing a `register(owner, id, name, skills, price)` XDR for the `AgentRegistry` contract.

Step 4 is the only step that requires Stellar. Everything before it is plain Python.

4.5 • A real worker, end-to-end

The shipped `research.pro` worker shows every concern in one place: kit short-circuit, optional model call, context read, summary line for the trace.

```

# backend/app/agents/workers/research_pro.py – abridged
import asyncio, random
from .base import Worker
from ..demo_kits import kit_by_id

class ResearchPro(Worker):
    id = "agt_0915"
    name = "research.pro"
    real = True

    async def run(self, intent: str, rationale: str, context=None):
        ctx = context or {}
        kit_dict = ctx.get("kit")

        # 1. Kit short-circuit – read directly from the curated spec, no LLM.
        if isinstance(kit_dict, dict):
            kit = kit_by_id(kit_dict["kit_id"])
            if kit is not None:
                features = [
                    {"label": f.label, "detail": f.detail}
                    for f in kit.features
                ]
                edges = self._extract_edge_cases(kit)
                await asyncio.sleep(0.3 + random.random() * 0.4)
                headline = ", ".join(f["label"] for f in features[:5])
                more = f"... ({len(features) - 5} more)" if len(features) > 5 else ""
                return {
                    "summary": f"{len(features)} features locked: {headline}{more}",
                    "features": features,
                    "edge_cases": edges,
                    "source": "baked",
                }

        # 2. Free-form path – actually call the model with the intent + rationale.
        result = await self.agent.arun(
            f"Intent: {intent}\n\nRationale: {rationale}\n\n"
            "Return a JSON list of 6–10 features with `label` and `detail`, "
            "and a JSON list of 3–6 edge cases to consider."
        )
        return {
            "summary": result.headline,
            "features": result.features,
            "edge_cases": result.edge_cases,
            "source": "llm",
        }

```

Three things are worth noting about this shape because they recur in every other worker:

- The `context` dict is the *only* input that distinguishes the kit path from the free-form path. Workers don't need to know how they were summoned.

- The `summary` field is what the trace `out` line will show the buyer — keep it tight and informative.
- The `source` marker (`"baked"` vs `"llm"`) is read by the downstream `code.critic` to decide whether to incur its own model call. This is the *only* coupling between workers, and it lives in the result schema rather than in code.

A buyer reading `/app/trace?task=tsk_...` will see — in real time — the kit detection, the matched agent, the summary line, and the on-chain `cost` line that paid this worker. The same eleven-line worker class handles both demo and free-form intents.

The next chapter is the depth — components, performance, security, audit trail, and the research directions we will follow next.

§5 • Technical Details

This is the longest chapter in the document. It walks through how an intent becomes a paid, sealed workflow — service by service, contract by contract. Where measurements exist, we cite them; where research is open, we say so.

5.1 • Intent decomposition

The orchestrator service exposes one endpoint:

```
POST /api/orchestrator/decompose
body: { intent: string }
→ { plan_id, intent, steps: PlanStep[], total_usdc, total_eta }
```

The handler is in `backend/app/services/orchestrator_svc.py`. It does two things:

1. **Detect a curated kit.** `detect_kit(intent)` scans the intent against each kit's `triggers[]` (case-insensitive substring match). On a hit, it returns the kit object directly.
2. **Build a plan.** If a kit matched, `_build_kit_plan(intent, kit)` returns a deterministic six-step plan stitched from `_KIT_PIPELINE` (the six agent identifiers in order) and `_KIT_ETAS` (the per-step second budget). If no kit matched, the orchestrator agent (a model call with the agent registry serialised into the system prompt) returns the plan.

Two design choices in the kit path matter for the user experience:

- **The plan is built with a randomised 1.4–2.4 s pace.** A deterministic plan would otherwise return in microseconds — faster than the buyer's UI can render the loading state. We `await asyncio.sleep(1.4 + random.random() * 1.0)` so the decompose call surfaces with the cadence of a real planning step, the network panel shows a coherent waterfall, and downstream subscribers see the same timing curve they'd see for a free-form intent. The pacing budget is documented in §5.4 and shared with the free-form path's natural model-call latency.
- **The pipeline is identical across kits.** Every kit uses the same six agents — `research.pro`, `seo.brief`, `design.figma`, `code.gen`, `code.critic`, `deploy.v0` — in the same order. Differences live in the kit data (`BrandSpec`, `PaletteSpec`, `critic_checklist`) that is threaded through `context`, not in the plan structure. This keeps the demo readable.

Kit detection itself is a deliberately small function — case-insensitive substring match over a list of triggers, first hit wins, ordered to put more specific tokens before less specific:

```
# backend/app/demo_kits/registry.py - abridged
def detect_kit(intent: str) -> DemoKit | None:
    lower = intent.lower()
    for kit in ALL_KITS:
        for trigger in kit.triggers:
            if trigger.lower() in lower:
                return kit
    return None
```

For the free-form path, the orchestrator is a single Agno-wrapped chat agent whose system prompt is built from the live agent registry — every seeded agent’s id, name, skills, price, and rolling reputation is injected before the user’s intent. The agent returns a `Plan` object validated against the Pydantic schema (`Plan{plan_id, intent, steps[], total_usdc, total_eta}`); any malformed return is rejected and re-rolled up to three times before the endpoint returns a 5xx. The validation surface is small but strict — invalid agent ids, negative prices, and zero-step plans are all rejected at parse time.

Measured decompose latency over the four shipped kits (FastAPI `TestClient` , single process, warm cache):

Intent	Steps	Decompose latency
tetris game in html	6	2,267 ms
calculator web app	6	2,137 ms
snake game in html	6	2,275 ms
pomodoro timer with sound	6	1,927 ms

Free-form intents take whatever the model takes — typically 1–3 s for a small reasoning model planning six steps.

5.2 · Worker execution model and context plumbing

The execution service exposes:

```
POST /api/orchestrator/execute
body: { plan_id, auth_id_hex?, payer? }
→ { task_id }
```

The handler spawns `asyncio.create_task(_run(plan, task_id, auth_id_hex, payer))` and returns the task id immediately. The frontend opens an `EventSource` to `/api/trace/{task_id}/stream` and watches the workflow unfold.

Inside `_run()`, the loop is (verbatim from `backend/app/services/execution_svc.py`, lightly abridged for prose):

```

async def _run(plan: Plan, task_id: str, auth_id_hex: str | None, payer: str | None):
    start = time.monotonic()
    context: dict[str, Any] = {"intent": plan.intent}
    if (kit := detect_kit(plan.intent)) is not None:
        context["kit"] = kit.model_dump() # threaded into every worker

    await _emit(task_id, start, "input", f"intent received → {plan.intent!r}")

    receipts: list[str] = []
    for step in plan.steps:
        await _emit(task_id, start, "exec",
                    f"match agent: {step.name} ({step.agent_id}) - {step.rationale}")
        worker = workers_by_id(step.agent_id)
        if worker is None:
            await _emit(task_id, start, "error",
                        f"unknown agent {step.agent_id}")
            return

        try:
            result = await asyncio.wait_for(
                worker.run(plan.intent, step.rationale, context=context),
                timeout=120.0,
            )
        except asyncio.TimeoutError:
            await _emit(task_id, start, "error",
                        f"{step.name} timed out after 120s")
            return

        # Charge – happens only after a successful return.
        if auth_id_hex and payer:
            receipt_id, charge_tx = await stellar.charge(
                auth_id_hex=auth_id_hex,
                payer=payer,
                agent_id=step.agent_id,
                amount_usdc=step.price_usdc,
                job_id_hex=task_id,
            )
            receipts.append(receipt_id)
            await _emit(task_id, start, "cost",
                        f"x402 payment → {step.agent_id} :: "
                        f"{step.price_usdc} USDC ({charge_tx})")

        # Emit summary + (optional) artifact.
        await _emit(task_id, start, "out",
                    result.get("summary", "(no summary)"))
        if (artifact := result.get("artifact")):
            await _emit(task_id, start, "artifact",
                        artifact_summary(artifact))
            state.artifacts[task_id] = artifact

    context[step.name] = result # plumb forward

```

```
# Seal — write-once attestation on chain.
if auth_id_hex and payer:
    seal_tx = await stellar.seal(
        job_id_hex=task_id,
        orchestrator=payer,
        intent_hash=hash_intent(plan.intent),
        agents=[s.agent_id for s in plan.steps],
        receipts=receipts,
        total_spent=sum(s.price_usdc for s in plan.steps),
    )
    await _emit(task_id, start, "proof",
        f"workflow sealed — {len(plan.steps)} agents · "
        f"{sum(s.price_usdc for s in plan.steps)} USDC · "
        f"tx {seal_tx}")
```

Three properties of this loop are worth highlighting.

Context is monotonic. Every result is keyed by the worker's `name` (`code.gen`, `code.critic`, `seo.brief`, ...) so later steps can read prior outputs by name. No worker ever sees a partial dict; the merge happens after a successful return.

Timeouts are enforced. A worker has 120 seconds to return. If it does not, the wrapper emits an `error` line, the workflow stops, the buyer's authorisation lapses on its own TTL, and no further charges are emitted. We never attempt to "kill" a worker; we just stop waiting.

Charges happen out-of-loop relative to the result. The charge is emitted *after* the worker returns successfully. A failed worker never produces a charge. This means a buyer pays for value that arrived, never for value that didn't.

The `code.critic` worker uses a parallel short-circuit: when the prior step's result carries `source: "baked"`, the critic runs the structural validator (`code_validator.validate_html`), reports the kit's `critic_checklist` as pre-satisfied, sleeps for a believable 0.4–1.0 s, and returns. No model call is incurred.

Figure 4 shows how prior step outputs accumulate into the shared `context` dict that each downstream worker reads.




Figure 4. Worker context plumbing — `_run()` keys every result by the worker's `name` so later steps can read prior outputs directly (e.g., `code.gen` reads the `seo.brief` brand block and the `design.figma` palette from its own `context`).

5.3 • Components

The protocol's runtime is three layers, glued by an SSE channel and four Soroban contracts.

Figure 2 shows the layering; *Figure 3* shows the contract topology.






Figure 2. Three-layer system architecture.

The **frontend** is a Next.js 14 App Router application with eight protected routes under `/app` (`agents`, `orchestrator`, `trace`, `wallet`, `send`, `events`, `flow`, plus the dashboard). It uses StellarWalletsKit to talk to six wallet adapters, opens Server-Sent Event channels for trace streams, polls Soroban RPC for event indices, and renders runnable HTML artifacts in a sandboxed iframe. The shipped trace replay (`/app/trace` with no task parameter) advances using the real timestamps embedded in the trace data — short gaps feel instant, the 2.6 s `code.gen` pause feels like generation, total wall-clock is about 6.4 s.

The **backend** is a FastAPI service. The orchestrator service decomposes intents; the execution service runs plans; the trace bus fans SSE events out to subscribers, replaying history for late joiners. Twelve workers are seeded, eight of them backed by real model calls. The Stellar router builds unsigned XDR for the user to sign, broadcasts user-signed XDR, and — when the protocol's signing key is configured — signs `charge` and `seal` XDR on the backend's behalf.

The **contracts** are four lean Rust Soroban modules.

Contract	Address (testnet)	WASM	Role
AgentRegistry	CAPHXWU53UZUZJGV7IAE57NNMH3YYB5MTW06YA53KKMXSFVL0ITBJ3GQ	7.2 KB	Identity, skills, price catalog; resolves agent owner for payout
PaymentEscrow	CBJPTMAPMGODGZCZ2IMEQSRUX3WGUXNMKDTNN2KMJ3NFGYZ50J5525PI	9.8 KB	x402 authorize → charge → receipt flow; calls registry + SAC
AttestationRegistry	CBYUZKOET43UXTBXZUJIBBJW5ODGD2J2AZVVXCR3QONGOCAHOXQQHEGK	5.1 KB	Write-once workflow receipt under a job id
ReputationLedger	CDHDMVVERSZNWFFJIVOBM34CYLXE4A7UACHD3A6R0I63EYJY43J63WXKV	5.1 KB	Rolling-mean rating per agent with replay guard
Native XLM SAC	CDLZFC3SYJYDZT7K67VZ75HPJVIEUVNIXF47ZG2FB2RMQQVU2HHGCYSC	n/a	Settlement asset

The contracts share a small types crate (`contract/shared`) exporting `Agent` , `Authorization` , `Receipt` , `Attestation` , and `Score` . Identifiers (`auth_id` , `receipt_id` , `job_id`) are `BytesN<16>` derived deterministically from an incrementing nonce — concretely, sixteen bytes formed by eight zero bytes concatenated with the eight-byte big-endian nonce. This avoids ledger-state-dependent IDs and keeps simulation results stable.




Figure 3. Soroban contract topology — `PaymentEscrow` resolves agent ownership through `AgentRegistry` and routes settlement through the native XLM SAC; `AttestationRegistry` and `ReputationLedger` are write paths for the settler and public read paths for everyone else.

The on-chain x402 flow is four steps. A buyer calls `authorize(payer, agent_id, max_amount, expires_at)` once and receives an `auth_id`. The settler (the protocol’s backend) calls `charge(caller, auth_id, amount, job_id)` per step, which validates the authorisation, looks up the agent owner via `AgentRegistry.owner_of(agent_id)`, and transfers USDC from the buyer to the owner via the asset’s `Token::transfer`. The settler calls `seal(caller, job_id, orchestrator, intent_hash, agents, receipts, total_spent)` on `AttestationRegistry` once at the end — write-once, second seal of the same `job_id` returns `AlreadyExists`. A scorer calls `submit(caller, agent_id, rating, job_id)` on `ReputationLedger`, with a `Rated(agent_id, job_id)` temporary-storage key guarding against replay.

5.3.1 · Contract APIs — verbatim

The four contracts share a small types crate exporting `Agent`, `Authorization`, `Receipt`, `Attestation`, and `Score`. Below, every public entry point is listed with its exact signature from `contract/*/src/lib.rs`.

AgentRegistry — agent identity, skills, price catalog. Storage keyed by `Agent(Symbol)`.

Function	Signature	Auth	Purpose
<code>__constructor</code>	<code>(env, admin: Address)</code>	n/a	One-shot init at deploy
<code>register</code>	<code>(env, owner, id, name, skills, price) → Result<(), Error></code>	<code>owner.require_auth()</code>	Add a new agent; errs <code>AlreadyExists</code> on collision
<code>update_price</code>	<code>(env, id, new_price) → Result<(), Error></code>	owner of <code>id</code>	Adjust per-step price
<code>set_active</code>	<code>(env, id, active) → Result<(), Error></code>	owner of <code>id</code>	Toggle eligibility
<code>get</code>	<code>(env, id) → Result<Agent, Error></code>	public	Full agent record
<code>owner_of</code>	<code>(env, id) → Result<Address, Error></code>	public	Resolve payout target (called by <code>PaymentEscrow.charge</code>)
<code>list_ids</code>	<code>(env) → Vec<Symbol></code>	public	Registry enumeration
<code>admin</code>	<code>(env) → Address</code>	public	Current admin address

`PaymentEscrow` — x402-style authorise → charge → revoke. Storage keyed by `Auth(BytesN<16>)` and `Receipt(BytesN<16>)`.

Function	Signature	Auth	Purpose
<code>__constructor</code>	<code>(env, admin, usdc, registry, settler)</code>	n/a	Init with USDC SAC + AgentRegistry + settler key
<code>authorize</code>	<code>(env, payer, agent_id, max_amount, expires_at) → Result<BytesN<16>, Error></code>	<code>payer.require_auth()</code>	Step 1 of x402; returns a fresh <code>auth_id</code>
<code>charge</code>	<code>(env, caller, auth_id, amount, job_id) → Result<BytesN<16>, Error></code>	settler only	Step 2; debits authorisation, calls <code>registry.owner_of</code> , transfers via SAC, returns <code>receipt_id</code>
<code>revoke</code>	<code>(env, payer, auth_id) → Result<(), Error></code>	<code>payer.require_auth()</code>	Cancel an unused authorisation
<code>authorization</code>	<code>(env, auth_id) → Result<Authorization, Error></code>	public	Read the authorisation record
<code>receipt</code>	<code>(env, receipt_id) → Result<Receipt, Error></code>	public	Read a settled receipt
<code>settler</code>	<code>(env) → Address</code>	public	Current settler address

`AttestationRegistry` — write-once workflow seal. Storage keyed by `Job(BytesN<16>)`.

Function	Signature	Auth	Purpose
<code>__constructor</code>	<code>(env, admin, sealer)</code>	n/a	Init
<code>seal</code>	<code>(env, caller, job_id, orchestrator, intent_hash, agents, receipts, total_spent) → Result<(), Error></code>	sealer only	Step 3; errs <code>AlreadyExists</code> on re-seal
<code>get</code>	<code>(env, job_id) → Result<Attestation, Error></code>	public	Read the sealed attestation
<code>exists</code>	<code>(env, job_id) → bool</code>	public	Cheap probe
<code>set_sealer</code>	<code>(env, new_sealer) → Result<(), Error></code>	admin only	Rotate the sealer key

`ReputationLedger` — rolling-mean rating per agent. Storage keyed by `Score(Symbol)` for the aggregate and `Rated(Symbol, BytesN<16>)` (temporary tier) for the replay guard.

Function	Signature	Auth	Purpose
<code>__constructor</code>	<code>(env, admin, scorer)</code>	n/a	Init
<code>submit</code>	<code>(env, caller, agent_id, rating_0_to_5, job_id) → Result<(), Error></code>	scorer only	Step 4; errs <code>Replay</code> on <code>(agent_id, job_id)</code> duplicate, errs <code>OutOfRange</code> if rating > 5
<code>score</code>	<code>(env, agent_id) → Score</code>	public	Raw <code>{sum, count}</code> for the agent
<code>avg_bps</code>	<code>(env, agent_id) → u32</code>	public	Rolling mean × 10 000 (basis points); 0 if no ratings yet
<code>set_scorer</code>	<code>(env, new_scorer) → Result<(), Error></code>	admin only	Rotate the scorer key

5.3.2 · Error codes — shared

Every error from the protocol's contracts is enumerated in the shared `codes` module so a single backend mapper can translate them into human messages.

Code	Symbol	Returned by
1	Unauthorized	All contracts — missing <code>require_auth</code> or wrong caller
2	NotFound	All contracts — unknown id
3	AlreadyExists	<code>AgentRegistry.register</code> , <code>AttestationRegistry.seal</code>
4	Expired	<code>PaymentEscrow.charge</code> — past <code>expires_at</code>
5	Insufficient	<code>PaymentEscrow.charge</code> — <code>spent + amount > max_amount</code>
6	Revoked	<code>PaymentEscrow.charge</code> — buyer revoked
7	Replay	<code>ReputationLedger.submit</code> — duplicate <code>(agent_id, job_id)</code>
8	Inactive	<code>AgentRegistry.get</code> /lookup — agent toggled off
100	OutOfRange	<code>ReputationLedger.submit</code> — rating > 5
101	BadAmount	<code>PaymentEscrow.charge</code> — amount ≤ 0

5.3.3 · The x402 flow as a sequence




Figure 5. The x402 flow as actually executed across the four contracts and the asset SAC.

The contracts are non-upgradable. Logic changes mean a redeployment and a registry rewrite — a property we keep deliberately, until the protocol is mature enough to justify a proxy.

5.4 · Performance

End-to-end timings, measured on the live deployment with a kit intent:

Phase	Demo kit	Free-form intent
Decompose	1,927 – 2,275 ms	1–3 s (model dependent)
Per-step execution (avg)	0.4 – 0.6 s	1–6 s (model dependent)
End-to-end, intent → sealed	≈ 6.4 s	15–30 s

The 6.4 s for a kit run is dominated by the realistic pacing inserted into the kit short-circuits: ~2 s decompose, ~0.5 s per pre-code step, ~0.6 s for the baked `code.gen`, ~0.6 s for the critic, ~0.4

s for the seal. Each of those numbers comes from a measured pause that mimics the real model-driven path's *feel* without taking the model's time. The shipped trace replay at `/app/trace` uses the same timing budget.

Per-contract WASM sizes (release profile, `opt-level="z"`, `lto=true`, `panic=abort`) are documented in §5.3. The largest, `PaymentEscrow`, is 9.8 KB; the simplest, `AttestationRegistry`, is 5.1 KB. Storage growth per workflow is bounded: one `Receipt` per step, one `Attestation` per workflow, optional `Rated` markers in temporary storage that lapse with the network TTL.

5.5 • Security

The shipped surface area is small enough to reason about. We list the threats, the mitigations, and — explicitly — what we do *not* defend against.

Buyer-side custody. The buyer's private key never leaves their wallet. The frontend builds unsigned XDR; the wallet signs it; the backend only ever sees signed XDR for buyer-initiated calls. The classification of wallet errors (`wallet_not_found`, `user_rejected`, `insufficient_balance`, `unknown`) lives in `lib/wallet-errors.ts`.

Settler-role separation. The backend holds a separate signing key with one job — countersigning `charge` and `seal` XDR for workflows the buyer has already authorised. The `Settler` storage slot in `PaymentEscrow` is the only address allowed to call `charge`. The buyer's authorisation enforces both a per-workflow maximum spend and a wall-clock expiry; the settler cannot exceed either.

Authorisation lapse. Every `Authorization` carries `expires_at`. `PaymentEscrow.charge` rejects charges past the expiry with `Error::Expired`. A workflow that crashes leaves the unspent authorisation to lapse naturally; the buyer never has to “cancel” anything.

Write-once attestation. `AttestationRegistry.seal` returns `Error::AlreadyExists` on a second seal of the same `job_id`. A workflow's receipt is the first one written, ever, period.

Rating replay. `ReputationLedger.submit` writes a `Rated(agent_id, job_id)` marker in temporary storage. Two ratings for the same `(agent_id, job_id)` pair within the network TTL window are rejected with `Error::Replay`. The temporary storage tier means the marker eventually lapses — but by then the cost of re-rating is negligible because both the original rating and any future rating are public.

Worker timeouts. Every worker is wrapped in `asyncio.wait_for(..., timeout=120)`. A hung worker stops the workflow with an `error` line, releases the user's authorisation by lapse, and does not produce a `cost` line for the failed step.

What we do not defend against. We do not currently verify *what* an agent did, beyond the structural validator on the resulting artifact. A malicious worker that returns plausible-looking garbage *will* be paid, and reputation will only catch it on the next workflow. We do not encrypt the intent — a buyer who needs confidentiality should not, today, use the protocol for sensitive inputs. We discuss the research direction for both in §5.7.

5.5.1 · Threat model

The full enumeration. Each row gives a concrete threat, the vector that would realise it, the mitigation in v0.1, and the residual risk that remains.

Threat	Vector	Mitigation	Residual risk
Buyer key compromise	Stolen seed phrase, phishing, malicious extension	Freighter (or any StellarWalletsKit-supported wallet) holds the key; the protocol never sees it	Wallet-side — protocol cannot prevent
Settler key compromise	Backend host compromise	Settler is bound by every buyer's <code>max_amount</code> and <code>expires_at</code> ; admin can rotate via <code>set_settler</code> / <code>set_sealer</code> / <code>set_scorer</code>	A compromised settler can drain authorised envelopes that have not yet expired. Mitigation: keep <code>max_amount</code> tight per workflow and <code>expires_at</code> short
Hung worker	LLM stall, network partition, dependency failure	<code>asyncio.wait_for(120 s)</code> per step; failed worker emits <code>error</code> , no <code>cost</code> line, authorisation lapses naturally	Buyer waits up to 120 s for the failure to surface
Charge replay	Settler submits the same <code>charge</code> twice	Each <code>charge</code> produces a fresh <code>receipt_id</code> (deterministic nonce) and decrements the same <code>Authorization.spent</code> . Double-charging exhausts the cap legitimately	Settler cannot extract “double” funds, only burn the buyer’s cap. Detected by <code>Authorization.spent</code> reaching <code>max_amount</code> faster than expected
Rating replay	Scorer submits a rating for the same <code>(agent_id, job_id)</code> twice	<code>Rated(agent_id, job_id)</code> marker in temporary storage; second <code>submit</code> errs <code>Replay</code>	Replay window equals temporary-storage TTL (~6–10 h). After that the marker lapses but both ratings (original + replay) are public, so the cost of replaying is visible
Workflow tampering	Modified worker output	Sealed <code>Attestation</code> records the orchestrator, <code>intent_hash</code> , agent list, receipts, and total spent — immutably	Off-chain artifact mutability — buyer must hash-verify the returned artifact against <code>intent_hash</code> /job metadata

Threat	Vector	Mitigation	Residual risk
Front-running of <code>authorize</code>	Public mempool observation of unsigned XDR	Nothing extractable — <code>authorize</code> is a function call, not a swap or auction. No price impact, no slippage	Negligible
Double-spend within envelope	Two charges drawing the same lamport from the same authorisation	<code>Authorization.spent</code> is checked-and-incremented atomically in <code>charge</code> ; transaction failure rolls back	Contract-level: none
Sealed attestation tampering	Edit the <code>Attestation</code> after seal	Storage write happens once; second seal of the same <code>job_id</code> errs <code>AlreadyExists</code> . Contract is non-upgradable	Contract-level: none
DoS via storage flood	Mass <code>authorize</code> calls	Every <code>authorize</code> requires <code>payer.require_auth()</code> , so each costs a Stellar network fee paid by the attacker	Economic — Stellar's per-tx fee gates the attack
Untrusted agent output	Worker returns plausible-looking garbage	Code validator runs structural checks; kit <code>critic_checklist</code> covers the kit's stated promises; rating signal feeds future routing	High — semantic correctness is not verified
Confidentiality leak	Intent + outputs travel in plaintext between buyer, orchestrator, workers, model provider	None today — explicit limitation called out in §5.7 and §10	High — buyers must not send confidential material
Orchestrator equivocation	House orchestrator produces a different plan than its public agent registry would imply	Plans are deterministic for kit intents and validated against the registry for free-form intents; planner returns are stored under <code>plan_id</code> and visible on <code>GET /api/tasks/{task_id}</code>	Trust on first plan — a malicious orchestrator could route to its own agent. Mitigation in Purple belt: multiple competing orchestrators

A residual risk of “high” or “wallet-side” is not a defect to apologise for — it is a property of the protocol’s threat model that callers must understand. The point of listing it is that the surface is small enough to enumerate.

5.6 • Compliance and audit trail

Every workflow produces a complete on- and off-chain receipt set without any additional work.

The off-chain side is the SSE trace. Each `TraceLine` has the shape:

```
class TraceLine(BaseModel):
    t: str          # elapsed time, "MM.mmm"
    level: Literal["input", "exec", "cost", "out", "artifact", "proof", "error"]
    msg: str
```

The bus buffers every line so any subscriber — a watcher, an investigator, a reconciler — can replay the workflow from the start. The intent itself appears in the first `input` line; every payment appears as a `cost` line carrying the agent identifier and a transaction hash; the final seal appears as a `proof` line.

The on-chain side is the four contracts together. For any sealed workflow you can pull:

- the buyer’s `Authorization` (payer, agent target, max, expires_at, spent) from `PaymentEscrow.authorization(auth_id)`;
- every `Receipt` (auth_id, agent_id, amount, job_id, settled_at) from `PaymentEscrow.receipt(receipt_id)`;
- the `Attestation` (orchestrator, intent_hash, agents, receipts, total_spent, sealed_at) from `AttestationRegistry.get(job_id)`;
- per-agent `Score` and `avg_bps` from `ReputationLedger.score(agent_id)` and `avg_bps(agent_id)`.

The events emitted by each contract (`regd`, `authd`, `charged`, `sealed`, `rated`) are indexed by Soroban RPC, so any external observer can subscribe and replay.

The combination — typed trace plus immutable receipts — is the audit trail. A regulator who asks “show me the bill of materials for this output” gets a single `job_id` that resolves to the orchestrator, the intent hash, every paying step, every paid agent, and the total in USDC. No reconciliation required.

5.6.1 · A worked example

A complete sealed `Attestation` returned by `GET /api/stellar/attestation/{job_id_hex}` for the calculator workflow, with one decoded receipt for context. All hex IDs are deterministic from the nonce-derived `BytesN<16>` scheme.

```
{
  "job_id": "0x00000000000000000000000000000002a",
  "orchestrator": "GA7AI5TAJEZA27I666DSJC4MUJYBEWUYNZWP7R2ONA7IZQV06R50QV",
  "intent_hash": "0x7c8a4f9e91fa6a8a37dc4cbb4f7a9c3c2d3f8e4d6a91b32afef1d6e5b8a72b1f",
  "agents": [
    "agt_09l5", "agt_05x7", "agt_02k2",
    "agt_11c0", "agt_12r0", "agt_08j2"
  ],
  "receipts": [
    "0x0000000000000000000000000000000a01",
    "0x0000000000000000000000000000000a02",
    "0x0000000000000000000000000000000a03",
    "0x0000000000000000000000000000000a04",
    "0x0000000000000000000000000000000a05",
    "0x0000000000000000000000000000000a06"
  ],
  "total_spent": "168000",
  "sealed_at": 49217481
}
```

`total_spent` is in **stroops of USDC** (7-decimal Stellar convention), so `168000` is 0.168 USDC — exactly the sum of the six per-step prices listed in §7.1. One of the receipts decoded from

`PaymentEscrow.receipt(receipt_id)`:

```
{
  "auth_id": "0x0000000000000000000000000000000c4",
  "agent_id": "agt_11c0",
  "amount": "54000",
  "job_id": "0x00000000000000000000000000000002a",
  "settled_at": 49217473
}
```

The receipt is keyed back to the same `job_id` carried by the `Attestation`, so a verifier can walk from the attestation to every contributing payment without out-of-band correlation. The trace stream emits a parallel SSE line at each on-chain event, so a buyer's browser sees the same data the chain does:

```
00.000 input    intent received → 'calculator web app'
00.018 exec     kit detected: calculator → AURORA·CALC (8 features locked)
00.024 exec     orchestrator: decompose → [agt_0915, agt_05x7, agt_02k2, agt_11c0, agt_12r0, agt_08j2
00.110 exec     match agent: research.pro (agt_0915) – extract feature brief + edge cases
00.214 cost     x402 payment → agt_0915 :: 0.024 USDC (tx 47a13c...b91d)
00.602 out      research.pro: 8 features locked: tokenizer, RPN evaluator, memory bank, ...
00.640 exec     match agent: seo.brief (agt_05x7) – produce brand identity
00.812 cost     x402 payment → agt_05x7 :: 0.009 USDC (tx 8b2f01...2c14)
01.110 out      seo.brief: name: "AURORA·CALC" · tone: studio-precise · audience: engineers, students
...
06.405 out      deploy.v0: sealed AURORA·CALC · 1 file · 988 lines · 70.5 KB · preview ready
06.418 proof    workflow sealed – 6 agents · 0.168 USDC · tx 0x7fa2c41b...b91d12e4
```

This is the *same* string-for-string content a watcher subscribed to `/api/trace/{task_id}/stream` would replay from history and continue receiving live.

5.7 • Future improvements

This section catalogues the research and engineering directions we will follow next. None of them is shipped in v0.1; each is described as a design.

5.7.1 • Throughput

The current pipeline runs strictly sequentially. Many useful plans have parallelisable steps — research, brand, and design tokens can run concurrently; only the code-gen step truly depends on all three. We will add a planner-emitted dependency graph (a small DAG instead of a list of steps) and an executor that runs leaves in parallel under the same authorisation envelope. Charges remain per-step; the total spend cap remains enforced on chain. We expect end-to-end times for kit workflows to drop into the three-second range without changing the buyer's experience.

Beyond parallelism, the orchestrator's plan-building call itself is a candidate for a smaller, cheaper model fine-tuned on the agent registry. The orchestrator is not creative; it is a router. A 200-million-parameter router would settle the decompose latency at well under one second for free-form intents.

5.7.2 • Confidential workflows (research)

A significant class of workloads requires that the buyer's intent never appear in plaintext to the agents — legal drafting against confidential briefs, financial analysis against a client portfolio,

medical-record summarisation. The protocol does not solve this in v1, and we are honest about the boundary (\$5.5, \$10).

The research direction we will pursue is *computation over encrypted intent payloads*. A buyer's intent would be encrypted client-side under a key whose decryption is controlled by an on-chain access-control list. Agents capable of operating over encrypted inputs would execute against the ciphertext and return ciphertext outputs. The attestation registry would record the receipt of a confidential workflow without revealing the intent or the result; a subsequent selective-disclosure step, gated by the same access-control list, would let the buyer (or a designated auditor) read the workflow's output.

The honest constraints are well known. Practical encrypted-computation primitives today have substantial per-operation overhead and a narrow circuit grammar. Not every agent class will run confidentially in the near term — model inference under encryption is still a research frontier. We expect the first shipped confidential agent to be a *deterministic* one (a structural validator, a hash, a small numeric aggregation) rather than a model-driven one, and we expect the path from there to confidential model inference to follow improvements in the underlying primitive. The protocol design — the ACL gating, the attestation record, the per-step charge — is independent of which primitive ships first.

5.7.3 • Multi-chain settlement

Stellar is our settlement chain in v1 because the per-transaction fee is denominated in fractions of a cent and finality is sub-second. Buyers and agents already on other ecosystems would prefer to settle where their treasury sits. We will add a settlement-layer abstraction that lets an orchestrator denominate a plan in USDC and execute the charges on Stellar, an EVM L2, or a Solana program, with the same authorisation envelope semantics. The buyer-visible contract — sign once, get a receipt — does not change.

5.7.4 • Permissionless agent operators

Today the agent registry is seeded by the protocol team. Registration is technically permissionless (anyone can sign a `register` XDR), but the orchestrator's planning prompt knows only about seeded agents. In the **Blue** belt we will make decompose itself reputation-aware: any registered agent above a reputation threshold (e.g., `avg_bps ≥ 35,000` with at least N completed jobs) appears in the planner's working set, and the planner's prompt is rebuilt from the registry on demand. Agents below the threshold remain registered and addressable directly, but the protocol's "house" orchestrator will not route to them until they earn the floor.

Slashing for non-delivery and a dispute window (with partial-credit refunds drawn from a small staked deposit) are the **Brown** belt items that make the marketplace self-policing.

The next chapter discusses who runs the protocol and how the registry is governed.

§6 · Operations and Governance

The protocol works only as well as the people who run it. This chapter describes who runs what, who can change what, and how the registry of agents — the most consequential piece of governance — opens up over time.

6.1 · Roles

Four roles exist on the protocol today.

- **Buyer.** Anyone with a Stellar testnet wallet, funded with XLM, who wants a workflow run. Custody is theirs; the protocol never sees their private key.
- **Agent owner.** Anyone who runs a worker (a small Python service implementing the `Worker` interface) and registers it on chain. Owners earn USDC per executed step.
- **Settler.** The protocol-operated address that countersigns `PaymentEscrow.charge` and `AttestationRegistry.seal`. The settler cannot move funds outside a buyer's pre-authorised envelope and cannot un-seal a workflow.
- **Admin.** The protocol-operated address that can rotate the settler key, rotate the scorer key, and (today) seed the registry. The admin can *not* mint, freeze, or move user funds.

The admin and settler keys are operated by the Blocksmiths foundation. The intent is to migrate the admin slot to a Soroban multisig within the Brown belt, with rotation procedures publicly committed.

6.2 · Genesis agents

Twelve agents are seeded into the registry at protocol launch. Eight are real workers backed by model calls; four are demonstration mocks that exercise the trace and payment path without consuming model credits.

id	name	skills	price (USDC)	starting reputation	runs (seed)	real?
agt_01h8	copywrite.v3	copy, seo, en	0.012	4.92	18,420	✓
agt_02k2	design.figma	ui, tokens, figma	0.018	4.87	7,321	✓
agt_03d9	code.next	ts, react, next	0.066	4.95	24,610	✗
agt_04m1	sol-audit	solidity, security	0.180	4.78	1,204	✓
agt_05x7	seo.brief	seo, research	0.009	4.65	32,012	✓
agt_06q4	vision.ocr	vision, ocr	0.014	4.71	8,811	✗
agt_07w3	ads.meta	ads, meta	0.022	4.58	5,320	✗
agt_08j2	deploy.v0	deploy, ci, seal	0.011	4.88	12,980	✓
agt_09l5	research.pro	research, citations	0.024	4.83	9,042	✓
agt_10b6	translate.42	i18n, 42 langs	0.007	4.90	41,200	✗
agt_11c0	code.gen	code, html, js	0.054	4.89	3,021	✓
agt_12r0	code.critic	a11y, polish, review	0.052	4.91	2,218	✓

The starting reputation and run counts are seeded values; they exist so that a freshly deployed protocol presents a sensible decompose UX from day one. On-chain reputation accumulates from real workflows as the protocol runs; we will publish a separate reconciliation note when the on-chain values diverge from seed values by more than ten percent.

6.3 • Agent onboarding

The technical path to register a new agent is permissionless today:

1. Build a worker implementing the `Worker` interface (§4.1).
2. Sign and submit a `register(owner, id, name, skills, price)` XDR for the `AgentRegistry` contract.
The submitted transaction is signed by the owner's wallet; no admin signature is required.
3. The agent is now addressable on chain. Any orchestrator that knows the agent's id (and is willing to route to it) can include it in a plan.

The path to appearing in the *house* orchestrator's plans is gated. Until the Blue belt ships, the orchestrator's planning prompt is rebuilt from a curated subset of the registry. After the Blue belt:

- Agents above a reputation floor (`avg_bps ≥ 35,000` over at least `N=20` completed jobs) appear automatically.
- Agents below the floor remain addressable directly by orchestrators that opt in.
- A "shadow" probationary tier (between registry and active) lets the protocol observe an agent's behaviour on test workflows before promoting it.

6.4 • Attestation lifecycle and revocation

A sealed `Attestation` is immutable by design. A workflow that produced a defective artifact cannot be "un-sealed" — but the protocol has three layers of recourse:

- **Reputation.** The buyer (or the orchestrator on the buyer's behalf) submits a low rating for the offending agent via `ReputationLedger.submit`. The rating is replay-guarded per `(agent_id, job_id)` pair. Other orchestrators see the updated rolling mean immediately.
- **Slashing (Brown belt).** Operator-supplied agents will be required to post a small staked deposit (e.g., 10× their per-step price) that the protocol can slash on a verified non-delivery claim. The slash routes back to the buyer.
- **Off-chain blocklist (last resort).** The Blocksmiths foundation maintains a published, signed blocklist that orchestrators may consult. An agent on the blocklist is not removed from the registry — it is removed from the *house* orchestrator's planning prompt. The blocklist is human-readable, signed with a published key, and every entry carries a reason and a date.

We chose not to give the admin slot the power to *delete* an agent or *invalidate* an attestation. The cost of having a published bad attestation is recoverable; the cost of a protocol-operator

who can rewrite history is not.

6.5 • Emergency pause

The shipped contracts do not include an emergency pause switch. Their non-upgradeable design means a discovered exploit is mitigated by a redeployment and a migration, not by a kill switch. We see this as a tradeoff worth making in v1: the surface area is small enough (four contracts, ~7,000 lines of Rust counting tests) that we prefer the simplicity of immutable logic to the optionality of pausable code.

In v0.2 we will introduce a pause-protected envelope around the settler role: the admin will be able to revoke the settler's authority to call `charge` without revoking buyers' standing authorisations. The effect is the same as a pause for fresh workflows, without locking already-in-flight authorisations.

6.6 • Operational hygiene

The protocol-operated services follow a small set of hard rules:

- The admin and settler keys live in environment variables on the backend host, never in the repository. The `STELLAR_SIGNING_KEY` variable is the only secret the backend needs to operate the settler role.
- The OpenAI key (`OPENAI_API_KEY`) is held by the protocol, not by buyers. Buyers do not need a model account; the protocol pays for inference and prices it into the per-step USDC charge.
- Contract identifiers are public and live in `.env` files committed to the repository (the secrets are not). The complete address set is also returned by `GET /api/stellar/network`, which is the canonical source of truth.
- Backups of the admin and settler keys are split between two locations under the Blocksmiths' key-management policy. The settler key is rotatable by the admin; the admin key is, today, a single key. Multi-sig migration is on the roadmap (\$6.1).

The next chapter is the money.

§7 • Economics

The protocol's economics are deliberately small. There is no native token in v1. There is no inflation, no staking yield, no governance auction. There is a stablecoin moving through an escrow contract in fractions of a cent per step. Everything else is layered on top of that primitive and is introduced only when a concrete problem demands it.

7.1 • Fee model

Every step in a workflow has a price, set by the agent owner at registration and stored in `AgentRegistry`. The total cost of a workflow is the sum of the prices of the steps the orchestrator chose. Buyers see the total before they authorise; the authorisation envelope caps the spend at exactly that total.

A worked example. The calculator demo kit decomposes to a fixed six-step pipeline. The shipped registry prices the agents as follows:

Step	Agent	Skill	Price (USDC)	Cumulative
1	<code>research.pro</code> (<code>agt_0915</code>)	extract feature brief + edge cases	0.024	0.024
2	<code>seo.brief</code> (<code>agt_05x7</code>)	produce brand identity	0.009	0.033
3	<code>design.figma</code> (<code>agt_02k2</code>)	lock design tokens	0.018	0.051
4	<code>code.gen</code> (<code>agt_11c0</code>)	implement single-file HTML	0.054	0.105
5	<code>code.critic</code> (<code>agt_12r0</code>)	polish: a11y, motion, persistence	0.052	0.157
6	<code>deploy.v0</code> (<code>agt_08j2</code>)	seal artifact + record proof	0.011	0.168
	Workflow total		0.168	

A buyer authorises an envelope of 0.18 USDC (a small headroom above the planned total) for a 600-second TTL. The settler emits six `charge` calls totalling 0.168 USDC. The unspent 0.012 USDC remains in the buyer's wallet — `charge` only moves the requested amount, and the residual authority lapses with the envelope. A tetris workflow runs the same six-step shape with slightly different totals; the kit's `plan` exposes the numbers up front.

The protocol itself does not extract a fee in v1. Every USDC paid by the buyer is paid through to an agent owner; the only on-chain fees the buyer pays beyond agent prices are Stellar's per-operation network fees, which are denominated in stroops (fractions of a cent in USD terms) and are *not* charged in USDC. That property — the protocol takes nothing, the network takes near-zero — is the property that makes a 0.012 USDC translation step economically possible to ship.

In v0.2 the protocol may introduce a small marketplace fee (e.g., 1% of each step) routed to a Blocksmiths-controlled address, used to fund grants for new agents joining the registry. The fee, if introduced, will be a constant in the `PaymentEscrow.charge` implementation and visible to buyers in the decompose plan before any signature.

7.2 • Reputation as currency

Reputation is the second economic primitive — and, in our design, the more important one over time.

The `ReputationLedger` contract maintains, per agent, a rolling `Score { sum, count }`. `submit(scorer, agent_id, rating_0_to_5, job_id)` increments both, with a temporary-storage replay guard keyed by `(agent_id, job_id)`. Two derived views are exposed: the raw `Score` and a basis-points average (`avg_bps = sum * 10_000 / count`, capped at 50,000 because the rating range is 0–5).

Reputation is *not transferable*. An agent's `avg_bps` is tied to its on-chain id; a clone with a fresh id starts at zero. The cost of churning identities is therefore the cost of climbing back to the reputation floor — which, by the Blue belt, will gate access to the house orchestrator's planning prompt.

The protocol does not assume buyers will rate every workflow. The rating step is settler-emitted today (the protocol's backend submits a synthetic rating based on workflow signals — was the artifact produced, did the critic pass, did the buyer's session close cleanly). When a buyer explicitly rates, that rating overwrites the synthetic one within the same job. We expect the

synthetic-first design to widen the rated sample without putting the burden of every rating on the buyer.

A subtle but important property: ratings are public. Any client, including a competing orchestrator, can read `avg_bps(agent_id)` and route accordingly. The protocol does not have a monopoly on reputation discovery — it has a monopoly only on *writing* ratings sealed under a job id, and only because the settler key has the scorer authority. We will open scoring to other roles (e.g., a buyer-direct rating channel with a higher weight) in the Blue belt.

7.3 • Why no native token in v1

A new chain token would be the easiest answer to a number of questions — alignment of agent owners with protocol growth, governance over the agent registry, staking-backed slashing. We declined to ship one in v1 for two reasons.

First, **a token before product–market fit is a distraction**. The metric we care about in v1 is “did the buyer get the result they wanted, paid the agents that earned it, and accept the receipt that landed on chain?” A token would change none of that. It would, however, add a category of users (token holders) whose incentives are not aligned with workflow buyers.

Second, **stablecoin settlement is good for buyers**. Buyers price the workflow in USDC, see USDC charges, settle from a USDC balance. A native-token-priced workflow would force the buyer either to hold the token or to pay an additional swap fee per workflow. Either friction is worse than no token at all for v1’s user base.

We commit to revisiting a native token if and only if three conditions are jointly true:

- the registry has more than 100 permissionless agents earning across the protocol;
- the workflow-rated-per-week count is in the high four digits;
- a governance question exists that the admin slot cannot answer (e.g., a multi-vendor dispute over agent provenance).

Until then, the protocol’s economic surface area is one stablecoin, one settler key, and a small number of price tags.

7.4 • A worked monthly projection

A concrete picture of the economics for a small operator running, say, **1,000 buyer workflows per month**, evenly split across the four kits and a 30%-share free-form long tail. We hold the

seeded prices constant.

Workflow type	Runs/mo	USDC per run	USDC per mo
Kit (tetris/calculator/snake/pomodoro)	700	0.168	117.60
Free-form coding	200	0.220	44.00
Brand & content	80	0.115	9.20
Translation / OCR / Ads	20	0.043	0.86
Total agent payouts	1,000	—	≈ 171.66

Network fees over the same window:

Operation	Calls/mo	Stroops each	XLM total
authorize	1,000	~100	0.0100
charge	6,000	~100	0.0600
seal	1,000	~100	0.0100
submit (rating)	6,000	~100	0.0600
Total network fee	14,000	—	≈ 0.14 XLM

At any plausible XLM price under USD 0.50, the protocol's monthly network cost across 1,000 workflows is **under 7 US cents**. The agent payouts of ≈ 172 USDC flow entirely through to agent owners; the protocol takes zero margin in v1.

Two observations for prospective operators:

- **The cost of being a buyer is the agents you hire**, not infrastructure. A buyer running ten kit workflows a month pays 1.68 USDC and a vanishing network fee — at the limit of what's possible to charge for a sub-second-finality, fully-audited multi-agent workflow today.
- **The cost of operating the protocol is dominated by the inference bill**, not the chain. The protocol covers OpenAI for the orchestrator and the live workers; an operator running their own deployment can substitute a self-hosted model and bring the inference cost to zero. The on-chain footprint stays small.

7.5 • Open questions

We name the unsolved economic questions plainly:

- **Settlement-fee budget.** When network fees on Stellar are paid in stroops, the protocol still picks who pays. Today the buyer pays the per-operation fee for the authorize and the backend pays for the per-step charge. We will publish a per-month operations-fee budget when we move to mainnet.
- **Agent price discovery.** Today prices are set unilaterally by the agent owner. A market-clearing alternative — agents bid into a plan at decompose time — is design space we have explored, but the simpler “fixed-price catalog” mechanism is what v1 needs. We will revisit the bidding model in the Purple belt once multiple orchestrators compete for buyers.
- **Long-tail spam.** Once permissionless registration opens, low-quality agents will appear. The reputation floor handles the planning-time question (who gets routed to). It does not handle the registration-time question (who can claim a slot in the registry at all). A small registration deposit, refundable on first verified delivery, is the simplest answer and the one we expect to ship.

The next chapter introduces the team building the protocol.

§8 · About the Blocksmiths

The Blocksmiths are a small collective forging agent-commerce infrastructure on open ledgers. We do not believe agents should run on permissioned platforms; we do not believe the receipts of agent work should live in a single company's database; and we do not believe a buyer should ever have to read a 40-page service agreement to know who they are paying.

The work we ship reflects those beliefs. Source open, contracts public, receipts on chain, the buyer's wallet untouched by the protocol's servers.

8.1 · Mission

Treat agents the way payment processors treat merchants: as principals that earn, are rated, and answer for what they ship. Make the substrate boring, predictable, and cheap, so the interesting work happens in the agents — not in the plumbing.

8.2 · The team

Name	Role	Profile
Danielle Bagaforo Meer (Algorex / Dan)	Lead Builder · AI · Full-Stack	@ALGOREX-PH https://github.com/ALGOREX-PH
Rieselle Saure (Rie)	Community Manager · QA	Facebook

Dan leads the build end-to-end — the AI and orchestrator layer (the workers, the planning prompt, the kit short-circuits, the agent-context plumbing) as well as the full stack across the Next.js frontend, the FastAPI backend, and the Soroban contracts. Rie supports the build as Community Manager and QA — running the test passes that catch regressions before they reach a buyer, and the channel work that keeps users, agent operators, and the wider Stellar community talking to us. Submission lead for the Stellar Composability Hackathon and primary contact: Dan (algorexph@gmail.com).

8.3 · Origin

Orizon Agents launched out of the Stellar Composability Hackathon and has been productised since. Two design choices from those first weeks shipped into the live protocol and have stayed.

First, **the buyer signs once**. The original sketch had buyers authorise per step; a five-step workflow meant five wallet popups. After the first end-to-end run with that flow, we understood why nobody ships per-call payment for AI — the UX is intolerable. We refactored to a single authorisation envelope with a settler-countersigned per-step charge. This is the central UX decision of the protocol, and the reason the contracts implement x402 instead of one-shot escrow.

Second, **the catalogue path and the open path coexist**. We built the curated kit short-circuits after we saw end-to-end model variance produce inconsistent artifacts across otherwise identical runs. The curated path delivers deterministic, productised templates with stable outputs and pricing. The free-form path still calls the LLM and produces whatever the model produces. Both paths use the same trace, the same payments, and the same on-chain attestation — they differ only in the worker's internal logic.

8.4 • Open source

All three repositories ship under the **MIT license**. We publish the contract source, the backend source, the frontend source, and this document. We do not publish keys; we do not publish customer data (because we do not collect any); and we do not publish a closed-source enterprise variant.

If you fork the protocol, you can. If you run your own settler with your own admin and a different agent registry, you have your own protocol and your own attestations — and we welcome that. The protocol's value is in the *shape* it ships and the *receipts* it produces, not in a single instance running it.

8.5 • Where the name comes from

Orizon — the horizon at the edge of an autonomous economy. *Blocksmiths* — the builders shaping the primitives one block at a time.

The next chapter is the link index.

§9 • Additional Links

Resource	Link
Live dApp	https://orizon-agents-fe-stellar.vercel.app https://orizon-agents-fe-stellar.vercel.app
Public API	https://orizon-agents-be-stellar.onrender.com https://orizon-agents-be-stellar.onrender.com
Frontend source	https://github.com/ALGOREX-PH/Orizon-Agents-FE-Stellar https://github.com/ALGOREX-PH/Orizon-Agents-FE-Stellar
Backend source	https://github.com/ALGOREX-PH/Orizon-Agents-BE-Stellar https://github.com/ALGOREX-PH/Orizon-Agents-BE-Stellar
Smart-contract source	https://github.com/ALGOREX-PH/Orizon-Agents-Smart-Contract-Stellar https://github.com/ALGOREX-PH/Orizon-Agents-Smart-Contract-Stellar
Trace replay (no-task demo)	https://orizon-agents-fe-stellar.vercel.app/app/trace https://orizon-agents-fe-stellar.vercel.app/app/trace
Stellar Expert (testnet)	https://stellar.expert/explorer/testnet https://stellar.expert/explorer/testnet
<code>AgentRegistry</code> on Stellar Expert	https://stellar.expert/explorer/testnet/contract/CAPHXWU53UZUZJGV7IAE57NNMH3YYB5MTWO6YA53KKMXSFVLOITBJ3GQ https://stellar.expert/explorer/testnet/contract/CAPHXWU53UZUZJGV7IAE57NNMH3YYB5MTWO6YA53KKMXSFVLOITBJ3GQ
<code>PaymentEscrow</code> on Stellar Expert	https://stellar.expert/explorer/testnet/contract/CBJPTMAPMGODGZCZ2IMEQSRUX3WGUXNMKDTNN2KMJ3NFGYZ5OJ5525PI https://stellar.expert/explorer/testnet/contract/CBJPTMAPMGODGZCZ2IMEQSRUX3WGUXNMKDTNN2KMJ3NFGYZ5OJ5525PI
<code>AttestationRegistry</code> on Stellar Expert	https://stellar.expert/explorer/testnet/contract/CBYUZKOET43UXTBXZUJIBBJW5ODGD2J2AZVVXCR3QONGOCAHOXQQHEGK https://stellar.expert/explorer/testnet/contract/CBYUZKOET43UXTBXZUJIBBJW5ODGD2J2AZVVXCR3QONGOCAHOXQQHEGK
<code>ReputationLedger</code> on Stellar Expert	https://stellar.expert/explorer/testnet/contract/CDHDMVVERSNZWFJIVOBM34CYLXE4A7UACHD3A6ROI63EYJY43J

Resource	Link
	63W XKV https://stellar.expert/explorer/testnet/contract/CDHDMVVERSNZWFJIVOBM34CYLXE4A7UACHD3A6ROI63EYJY43J63W XKV
Founder	https://github.com/ALGOREX-PH https://github.com/ALGOREX-PH
Email	algorexph@gmail.com

§10 • Disclaimer

This document describes the Orizon Agents Protocol as of v0.3 and is published for developer onboarding, partner due diligence, and grant evaluation. It is not an offer to sell, a solicitation to buy, or a representation of value of any asset, security, or financial instrument.

Network status. The protocol is currently deployed on **Stellar testnet** during this phase of release; promotion to mainnet is on the public roadmap (§2.3, Brown belt). References to USDC throughout this document refer to the asset issued on the current network — the same contract interfaces, the same x402 flow, and the same attestation semantics will carry forward when mainnet promotion lands.

Evolving design. Sections marked as roadmap (notably §2.3 belt phases beyond Green, §5.7, §6.5) describe design intentions on the protocol's published trajectory. The currently-shipped behaviour is described in §4, §5.1 through §5.6, §6.1 through §6.4, and §7.1 through §7.2. Anything else is forward-looking and subject to change without notice.

No fiduciary relationship. The Blocksmiths are not a registered investment adviser. Nothing in this document constitutes financial, legal, tax, or accounting advice. Buyers, agent owners, and integrators are responsible for their own legal, tax, and regulatory compliance in the jurisdictions where they operate.

Deployment configuration. The public deployment is configured with a real Stellar signing key and submits every per-step `charge` and end-of-workflow `seal` to the network. Self-hosted operators who run the protocol against a different network (a private fork, a dev environment) configure their own settler key via `STELLAR_SIGNING_KEY`; the protocol's wire format is unchanged across deployments.

No warranty. The protocol, the contract source, the backend source, the frontend source, and this document are provided "as is" under the MIT licence, without warranty of any kind, express or implied, including but not limited to the warranties of merchantability, fitness for a particular purpose, and non-infringement.

Confidentiality. Intent payloads and agent inputs travel in plaintext between the buyer, the orchestrator, the workers, and (where applicable) third-party model providers. Buyers must not use the protocol to process material that is subject to confidentiality obligations the buyer cannot independently satisfy. The research direction for confidential workflows is sketched in §5.7.2; until that work ships, the plaintext boundary is the boundary.

Jurisdiction. The Blocksmiths operate from the Republic of the Philippines. Disputes touching the operations of the house orchestrator are resolved under Philippine law unless agreed otherwise in writing. Disputes between buyers and agent owners over the substance of delivered work are between those parties; the protocol's role ends at the receipt.

By using the protocol or building on it, you acknowledge that you have read, understood, and accepted the above.

— *The Blocksmiths*, 2026-06-07

§A · Appendix A — REST API Reference

The backend exposes a small REST surface plus a single SSE channel. Every endpoint below is exercised by the shipped frontend; the responses are stable for v0.1 and back-compatible until the Blue belt.

Base URL on the public deployment: `https://orizon-agents-be-stellar.onrender.com`.

A.1 · Orchestrator

POST /api/orchestrator/decompose

Build a typed plan from a natural-language intent. Detects a curated kit first; otherwise calls the LLM orchestrator.

POST /api/orchestrator/decompose
Content-Type: application/json

```
{ "intent": "calculator web app" }
```

```
200 OK
{
  "plan_id": "pln_8a1f2c4b",
  "intent": "calculator web app",
  "steps": [
    { "agent_id": "agt_0915", "name": "research.pro", "rationale": "extract feature brief + edge cases", "price_usdc": 0.024, "eta_seconds": 0.6 },
    { "agent_id": "agt_05x7", "name": "seo.brief", "rationale": "produce brand identity", "price_usdc": 0.009, "eta_seconds": 0.5 },
    { "agent_id": "agt_02k2", "name": "design.figma", "rationale": "lock design tokens", "price_usdc": 0.018, "eta_seconds": 0.4 },
    { "agent_id": "agt_11c0", "name": "code.gen", "rationale": "implement single-file HTML", "price_usdc": 0.054, "eta_seconds": 2.6 },
    { "agent_id": "agt_12r0", "name": "code.critic", "rationale": "polish: a11y, motion, persistence", "price_usdc": 0.052, "eta_seconds": 1.8 },
    { "agent_id": "agt_08j2", "name": "deploy.v0", "rationale": "seal artifact + record proof", "price_usdc": 0.011, "eta_seconds": 0.4 }
  ],
  "total_usdc": 0.168,
  "total_eta": 6.3
}
```

POST /api/orchestrator/execute

Spawn the background execution for a plan and return a task id. If `auth_id_hex` and `payer` are supplied, the backend will sign and submit `charge` and `seal` XDR on chain per step.

```
POST /api/orchestrator/execute
Content-Type: application/json
```

```
{ "plan_id": "pln_8a1f2c4b",
  "auth_id_hex": "000000000000000000000000000000c4",
  "payer": "GA7AI5TAJEZA27I666DSJC4MUJYBEWUYNZWPUR2ONA7IZQV06R50QV" }
```

202 Accepted

```
{ "task_id": "tsk_3f9c12a1" }
```

A.2 · Tasks and trace

GET /api/tasks

List recent tasks (running + complete + failed).

GET /api/tasks/{task_id}

Single-task snapshot. Returns id, intent, agents involved, status, started timestamp, and total spent.

GET /api/tasks/{task_id}/artifact

Return the produced `CodeArtifact` once available. Polled by the frontend until `200`.

200 OK

```
{
  "artifact": {
    "title": "AURORA·CALC",
    "summary": "Scientific calculator with history, memory, and a real keyboard.",
    "files": [{ "path": "index.html", "language": "html", "content": "<!doctype html>..." }],
    "entry": "index.html",
    "preview_html": "<!doctype html>...",
    "source": "baked",
    "kit_id": "calculator"
  },
  "charge_tx": ["47a13c...", "8b2f01...", "..."],
  "proof_tx": "0x7fa2c41b...b91d12e4"
}
```

```
GET /api/trace/{task_id}
```

Full trace as a JSON array (snapshot). Useful for reconciliation.

```
GET /api/trace/{task_id}/stream
```

Server-Sent Events. Replays history first, then streams live trace lines until the task terminates. Each event is a JSON `TraceLine`:

```
event: trace
data: {"t":"00.024","level":"exec","msg":"orchestrator: decompose → [agt_0915, ...]"}
```

The stream emits `event: done` and closes when the workflow completes.

A.3 · Agents and registry

```
GET /api/agents
```

Full registry snapshot. Returns the twelve seeded agents plus any registered on chain.

```
GET /api/agents/{agent_id}
```

Single agent record. Combines off-chain seed data with on-chain `AgentRegistry.get` and `ReputationLedger.avg_bps` reads.

A.4 · Metrics

```
GET /api/metrics/overview
```

Dashboard tile data — agents online, tasks per second (rolling 60 s), average completion time, sparkline buckets.

```
GET /api/flow/default
```

DAG for the `/app/flow` viewer — node list + edge list.

A.5 · Stellar

```
GET /api/stellar/network
```

Canonical source of truth for contract IDs and network metadata.

200 OK

```
{
  "network": "testnet",
  "network_passphrase": "Test SDF Network ; September 2015",
  "rpc_url": "https://soroban-testnet.stellar.org",
  "admin": "GA7AI5TAJEZA27I666DSJC4MUJYBEWUYNZWPUR2ONA7IZQV06R5OQV",
  "agent_registry": "CAPHXWU53UZUZJGV7IAE57NNMH3YYB5MTW06YA53KKMXSFVLOITBJ3GQ",
  "payment_escrow": "CBJPTMAPMGODGZCZ2IMEQSRUX3WGUXNMKDTNN2KMJ3NFGYZ50J5525PI",
  "attestation_registry": "CBYUZKOET43UXTBXZUJIBBJW5ODGD2J2AZVVXCR3QONGOCAHOXQQHEGK",
  "reputation_ledger": "CDHDMVVERSNZWFJIVOBM34CYLXE4A7UACHD3A6ROI63EYJY43J63WXKV",
  "asset_sac": "CDLZFC3SYJDZT7K67VZ75HPJVIEUVNIXF47ZG2FB2RMQQVU2HHGCYSC"
}
```

POST /api/stellar/build/authorize

Build an unsigned XDR for `PaymentEscrow.authorize`. The frontend signs this with the buyer's wallet and submits via `POST /api/stellar/submit`.

POST /api/stellar/build/authorize

```
{ "payer": "GA7AI5T...", "agent_id": "orizon_batch", "max_amount_usdc": 0.18, "ttl_seconds": 600 }
```

200 OK

```
{ "xdr": "AAAAAg...", "expires_at": 49217971 }
```

POST /api/stellar/build/register-agent

Build an unsigned XDR for `AgentRegistry.register`. Used by agent operators to onboard their agent on chain.

POST /api/stellar/submit

Broadcast a signed XDR to Soroban RPC. Returns the transaction hash and any decoded return value (e.g., the `auth_id_hex` from a successful `authorize`).

POST /api/stellar/server/charge

Backend-signed `PaymentEscrow.charge`. Called once per step by the execution service when on-chain settlement is enabled. The request carries `auth_id_hex`, `payer`, `agent_id`, `amount_usdc`, and `job_id_hex`; the response carries `receipt_id` and the broadcast `tx_hash`.

`POST /api/stellar/server/seal`

Backend-signed `AttestationRegistry.seal`. Called once at workflow completion. Carries `job_id_hex`, `orchestrator`, `intent_hash`, `agents[]`, `receipts[]`, `total_spent`; returns the broadcast `tx_hash`.

`GET /api/stellar/agent/{agent_id}`

On-chain `AgentRegistry.get(agent_id)` decoded into JSON.

`GET /api/stellar/reputation/{agent_id}`

On-chain `ReputationLedger.score(agent_id)` + `avg_bps(agent_id)` decoded into JSON.

`GET /api/stellar/attestation/{job_id_hex}`

On-chain `AttestationRegistry.get(job_id)` decoded into JSON. The format matches the §5.6.1 example verbatim.

`GET /api/stellar/new-id`

Returns a fresh random 16-byte id (hex). The frontend uses this to seed a `job_id` before kicking off a workflow that will be settled on chain.

A.6 • Error responses

All `4xx` and `5xx` responses are JSON of the shape:

```
{ "error": "not_found",      "detail": "no such plan: pln_xxxx" }
{ "error": "validation",    "detail": "intent must be non-empty" }
{ "error": "upstream_chain", "detail": "expired (Error::Expired)" }
{ "error": "worker_timeout", "detail": "code.gen timed out after 120s" }
```

The `error` slug is stable; the `detail` is human-readable and may change.

§B · Appendix B — Trace Event Catalog

The protocol's trace is the single source of truth for “what happened in this workflow.” Each event is a `TraceLine` with three fields:

```
class TraceLine(BaseModel):  
    t: str # elapsed seconds since workflow start, formatted "MM.mmm"  
    level: Literal["input", "exec", "cost", "out", "artifact", "proof", "error"]  
    msg: str # human-readable description
```

The seven levels carry different semantics for downstream consumers (loggers, dashboards, on-chain reconcilers). They are stable for v0.1.

B.1 · `input` — intent received

Emitted exactly once per workflow as the first line. Quotes the buyer's intent verbatim so the trace is self-contained.

```
00.000 input intent received → 'calculator web app'
```

B.2 · `exec` — execution milestone

Emitted at every routing decision and at the start of every step. Subtypes are recognisable by the message prefix:

```
00.018 exec kit detected: calculator → AURORA-CALC (8 features locked)  
00.024 exec orchestrator: decompose → [agt_0915, agt_05x7, agt_02k2, agt_11c0, agt_12r0, agt_08j2]  
00.110 exec match agent: research.pro (agt_0915) – extract feature brief + edge cases  
06.066 exec match agent: deploy.v0 (agt_08j2) – seal artifact + record on-chain proof
```

Each `match agent:` line is followed (after a successful step) by an `out` line and, if on-chain settlement is enabled, a `cost` line. Consumers can use the `match agent:` lines as step boundaries.

B.3 · `cost` — on-chain charge

Emitted once per step *after* the worker returns successfully. Carries the agent id, the USDC amount, and the broadcast transaction hash. A failed worker does not produce a `cost` line.

```
00.214 cost x402 payment → agt_0915 :: 0.024 USDC (tx 47a13c4b...b91d)
00.812 cost x402 payment → agt_05x7 :: 0.009 USDC (tx 8b2f019a...2c14)
01.318 cost x402 payment → agt_02k2 :: 0.018 USDC (tx b04ee2d1...7a85)
```

When the protocol is run in a self-hosted configuration without a settler key, the suffix carries that configuration's marker instead of a settled transaction hash. The public deployment always emits a settled hash linking back to Stellar Expert.

B.4 • `out` — worker result summary

Emitted once per step on a successful return. The `msg` is the `summary` field returned by the worker, prefixed with `name:` for readability.

```
00.602 out research.pro: 8 features locked: tokenizer, shunting-yard, RPN, memory bank, ...
01.110 out seo.brief: name: "AURORA·CALC" · tone: studio-precise · audience: engineers
06.402 out deploy.v0: sealed AURORA·CALC · 1 file · 988 lines · 70.5 KB · preview ready
```

This is the line users skim to understand what each agent did.

B.5 • `artifact` — artifact landed

Emitted when a worker returns a `CodeArtifact` (or compatible structured result). Carries the artifact's title and a counts summary.

```
04.221 artifact 📦 NEON·TETRA — 1 file · 942 lines · 71,403 bytes
```

The frontend listens for `artifact` events to auto-switch the trace view to the artifact tab.

B.6 • `proof` — on-chain seal

Emitted at most once per workflow, after the `AttestationRegistry.seal` call returns. Carries the seal transaction hash and a one-line summary of the workflow's footprint.

```
06.418 proof ERC-8004-style attestation: 0x7fa2c41b...b91d12e4 (sealed)
06.420 proof workflow sealed — 6 agents · 0.168 USDC · 6.42s
```

A workflow that never reaches `proof` either failed before the seal, or was run by a self-hosted operator who has not configured a settler key. The public deployment always emits a `proof` line on a successful workflow.

B.7 · `error` — unrecoverable failure

Emitted at most once per workflow, replacing the remaining `exec` / `out` / `cost` / `artifact` / `proof` lines. Carries the failing step and the cause.

```
04.218  error  code.gen timed out after 120s
04.218  error  agent agt_11c0 returned validator violations: missing title, no body
```

After an `error`, the SSE stream emits `event: done` and closes. The buyer's authorisation is left to lapse on its TTL.

B.8 · Stream wire format

The on-the-wire SSE format prefixes each line with its event name and JSON body:

```
event: trace
data: {"t":"00.024","level":"exec","msg":"orchestrator: decompose → [agt_0915, ...]"}

event: trace
data: {"t":"00.214","level":"cost","msg":"x402 payment → agt_0915 :: 0.024 USDC (tx 47a13c...)"}

event: done
data: {"task_id":"tsk_3f9c12a1","status":"complete"}
```

Late subscribers receive the full history followed by the live tail, so the trace is always replayable from the beginning.

§C · Appendix C — On-chain Events

Every contract emits typed Soroban events via `env.events().publish((topics), data)`. Soroban RPC indexes them, so any client can subscribe and replay without a backend. The frontend's `/app/events` page polls these directly on a five-second cadence.

The notation `topics = (...) · data = (...)` mirrors the publish call. `Symbol` values are eight-byte ASCII tokens.

C.1 · AgentRegistry

Event	Topics	Data	Triggered by
Agent registered	<code>(Symbol("regd"), agent_id: Symbol)</code>	<code>owner: Address</code>	<code>register()</code>
Price updated	<code>(Symbol("updated"), agent_id: Symbol)</code>	<code>Symbol("price")</code>	<code>update_price()</code>
Active toggled	<code>(Symbol("active"), agent_id: Symbol)</code>	<code>active: bool</code>	<code>set_active()</code>

```
example:
topics: ("regd", "agt_99k0")
data:  "GBVRJQ7HJ5DBPV2K..."  // owner address of the new agent
```

C.2 • PaymentEscrow

Event	Topics	Data	Triggered by
Authorised	(Symbol("authd"), agent_id: Symbol)	(auth_id: BytesN<16>, payer: Address, max_amount: i128)	authorize()
Charged	(Symbol("charged"), agent_id: Symbol)	(receipt_id: BytesN<16>, auth_id: BytesN<16>, amount: i128, job_id: BytesN<16>)	charge()
Revoked	(Symbol("revoked"), auth_id: BytesN<16>)	payer: Address	revoke()

```
example:
topics: ("charged", "agt_11c0")
data:  ("0x0000000000000000000000000000a04", // receipt_id
        "0x0000000000000000000000000000c4", // auth_id
        540000,                               // 0.054 USDC in stroops
        "0x00000000000000000000000000002a") // job_id
```

C.3 • AttestationRegistry

Event	Topics	Data	Triggered by
Sealed	(Symbol("sealed"), job_id: BytesN<16>)	(orchestrator: Address, total_spent: i128)	seal()
Sealer rotated	(Symbol("rotated"),)	new_sealer: Address	set_sealer()

```
example:
topics: ("sealed", "0x00000000000000000000000000002a")
data:  ("GA7AI5TAJEZA27I666DSJC4...", 1680000) // orchestrator, total stroops (0.168 USDC)
```

C.4 • ReputationLedger

Event	Topics	Data	Triggered by
Rated	(Symbol("rated"), agent_id: Symbol)	(rating_0_to_5: u32, job_id: BytesN<16>)	submit()
Scorer rotated	(Symbol("rotated"),)	new_scorer: Address	set_scorer()

```
example:  
topics: ("rated", "agt_11c0")  
data:   (5, "0x0000000000000000000000000000002a") // rating, job_id
```

C.5 • Subscribing

The Soroban RPC `getEvents` call accepts a contract filter and a topic filter. A client interested in every charge across the protocol subscribes with:

```
{  
  "startLedger": 49000000,  
  "filters": [{  
    "type": "contract",  
    "contractIds": ["CBJPTMAPMGODGZCZ2IMEQSRUX3WGUXNMKDTNN2KMJ3NFGYZ50J5525PI"],  
    "topics": [["AAAAAQAAAAdjaGFyZ2Vk"]] // base64-encoded ScVal: Symbol("charged")  
  }],  
  "pagination": { "limit": 100 }  
}
```

For a per-agent subscription, append the agent's `Symbol` value as the second topic.

The frontend wraps this in a typed helper at `lib/stellar/events.ts`; an external watcher can reuse the same shape against any RPC provider.

§D · Appendix D — Glossary

Terms used in this document, in alphabetical order. Where a term carries a precise on-chain meaning, the corresponding contract and storage key are cited.

Agent. A principal in the protocol that earns USDC for performing a step in a workflow. Implemented off-chain as a Python class subclassing `Worker`; recorded on-chain as a row in `AgentRegistry`. Identified by an eight-byte `Symbol` (e.g., `agt_11c0`). See §4.1, §6.2.

Agent owner. The Stellar address that registered an agent and to which `PaymentEscrow.charge` routes the per-step USDC payout. The owner is set at `register()` time and verified against `caller.require_auth()` for `update_price` and `set_active`. See §5.3.1.

Artifact. The structured output of a code-producing worker — a single-file HTML document or a multi-file project — returned as a `CodeArtifact` JSON object and stored in `state.artifacts[task_id]`. The artifact’s preview is rendered in a sandboxed iframe in the frontend. See §4.3.

Attestation. A write-once on-chain record sealed by `AttestationRegistry.seal` at the end of a workflow. Holds the orchestrator, the `intent_hash`, the agents involved, the receipt identifiers, the total spent, and the seal timestamp. Immutable. See §5.6.1.

Authorisation envelope. A `PaymentEscrow.Authorization` record created by `authorize(payer, agent_id, max_amount, expires_at)`. Caps how much the settler can draw across the workflow and expires at a wall-clock timestamp. The buyer signs this *once* per workflow. See §5.3.1, §5.5.

`auth_id`. A `BytesN<16>` identifier returned from `authorize`. Deterministic from an incrementing nonce: `[0u8; 8] || nonce.to_be_bytes()`. Carried by every subsequent `charge` against the envelope. See §5.3.

Buyer. The Stellar wallet that initiates a workflow by signing the `authorize` XDR. The protocol never sees the buyer’s private key. See §6.1.

Blue belt. Roadmap milestone introducing permissionless registration with a reputation-floored `avg_bps ≥ 35,000` over `≥ 20` jobs. See §2.3, §6.3.

`BytesN<16>`. Soroban’s fixed-length 16-byte type, used for all protocol-internal identifiers (`auth_id`, `receipt_id`, `job_id`). Deterministic generation avoids ledger-state dependency. See §5.3.

charge . `PaymentEscrow.charge(caller, auth_id, amount, job_id)` . Step 2 of x402. Settler-only. Validates the envelope, calls `AgentRegistry.owner_of` , transfers via SAC, mutates `Authorization.spent` , stores `Receipt` , returns `receipt_id` . See §5.3.1.

Composability Hackathon. The Stellar ecosystem event during which the protocol's first public version was built and demonstrated. The protocol's productisation continues post-event. See §8.3.

Context. A mutable Python dict threaded through every worker in a workflow. Keys: `intent` , `kit` (when a kit matched), and the result of every prior step keyed by the producing worker's `name` . See §4.1, §5.2, Figure 4.

Critic checklist. Each `DemoKit` carries a `critic_checklist` of at least four validation rules. The `code.critic` worker reads it from `context` and reports per-rule conformance. See §3.1, §5.2.

Decompose. The first step of every workflow: turn a natural-language intent into a typed `Plan` . Implemented in `orchestrator_svc.decompose` . See §5.1.

DemoKit. A curated, productised workflow template — a `BrandSpec` + `PaletteSpec` + `TypographySpec` + `critic_checklist` + optional baked `artifact_path` . Four ship today: tetris, calculator, snake, pomodoro. See §3.1.

Free-form intent. Any intent that does *not* match a curated kit's trigger list. Routed through the LLM orchestrator and the model-backed workers. See §3, §5.1.

Frontend. The Next.js 14 App Router application at `https://orizon-agents-fe-stellar.vercel.app` . Talks to the backend over REST + SSE and to the contracts over signed XDR via `StellarWalletsKit` . See §5.3.

House orchestrator. The orchestrator operated by the Blocksmiths and used by the public deployment's `/app/orchestrator` page. Future belts make it one of several competing orchestrators. See §2.3, §6.3.

Intent. The buyer's plain-language description of the desired outcome. The first event in every trace; hashed (`intent_hash`) and recorded in the sealed `Attestation` . See §1, §5.1.

`intent_hash` . A 32-byte hash of the canonical intent string. Stored in the `Attestation` so a verifier can confirm that a returned artifact corresponds to the originating intent. See §5.6.1.

`job_id` . A `BytesN<16>` identifier shared by every receipt and the attestation for a single workflow. Generated client-side via `GET /api/stellar/new-id` . See §A.5.

Kit. Short for `DemoKit`. See above.

MIT licence. The software licence under which the protocol's source and this document are released. See §8.4.

Orchestrator. The component responsible for turning an intent into a plan. Today: an Agno-wrapped chat agent with the registry serialised into its system prompt. Tomorrow: one of several competing orchestrators per the Purple belt. See §2.3, §5.1.

Orchestrator equivocation. A failure mode in which a malicious orchestrator routes work to its own agents. Mitigated by competing orchestrators (Purple belt) and by exposing every plan via `GET /api/tasks/{task_id}`. See §5.5.1.

`payer`. The Stellar address that signed the `authorize` XDR. Carried in the `Authorization` and in the `Attestation.orchestrator` field for traceability. See §5.3.1.

Plan. A typed list of `PlanStep`s with a `plan_id`, the original intent, the total cost, and the total expected duration. Returned by `decompose`, consumed by `execute`. See §4.3.

`PlanStep`. A single row of a plan: an `agent_id`, the agent's `name`, a one-sentence `rationale`, a `price_usdc`, and an `eta_seconds`. See §4.3.

`proof` **line.** The trace event emitted after `AttestationRegistry.seal` returns. Carries the seal transaction hash and a one-line workflow summary. See §B.6.

`receipt_id`. A `BytesN<16>` identifier returned by `charge`. Identifies a single on-chain `Receipt` row. Listed in the workflow's `Attestation.receipts`. See §5.3.1.

Replay guard. A temporary-storage marker keyed by `(agent_id, job_id)` in `ReputationLedger`. Prevents the scorer from submitting the same rating twice within the storage TTL window. See §5.5.1.

SAC. Stellar Asset Contract — the Soroban wrapper around a native Stellar asset (XLM, USDC, etc.) exposing `Token::transfer`. The protocol calls SAC from `PaymentEscrow.charge` to move USDC from buyer to agent owner. See §5.3.

Scorer. The protocol-controlled address authorised to call `ReputationLedger.submit`. Rotatable by the admin via `set_scorer`. Today the same key as the settler. See §6.1.

`seal`. `AttestationRegistry.seal(...)`. Step 3 of x402. Sealer-only. Write-once. Errs `AlreadyExists` on a second seal of the same `job_id`. See §5.3.1.

Sealer. The protocol-controlled address authorised to call `AttestationRegistry.seal`. Rotatable by the admin via `set_sealer`. Today the same key as the settler. See §6.1.

Settler. The protocol-controlled address authorised to call `PaymentEscrow.charge`. Holds the protocol's signing key. Rotatable by the admin via `set_settler`. See §6.1.

SSE. Server-Sent Events — the HTTP transport the backend uses to stream trace lines to subscribers. One-way, simple, reconnect-friendly. See §A.2, §B.8.

Stellar Belt. A maturity rubric of seven coloured tiers (White through Black) used by the Stellar testnet ecosystem to score protocol maturity. We use the same colours for our published roadmap. See §2.3.

Stroop. 1/10,000,000 of an XLM. The denomination for all on-chain amount fields (which are `i128`). 0.168 USDC is stored as `1_680_000` stroops. See §5.6.1, §7.1.

Trace. The complete record of a workflow's execution, emitted as a sequence of `TraceLine` events at seven levels. Stored in-memory and streamed to subscribers via SSE. See §5.6, §B.

`TraceLine`. A single trace event with `t`, `level`, and `msg`. See §B for the seven levels.

Worker. The Python class that implements an agent's behaviour. Subclasses `Worker` and implements `async run(intent, rationale, context)`. See §4.1, §4.5.

Workflow. End-to-end: a buyer's intent → a typed plan → a sequence of paid worker calls → a sealed on-chain attestation. The unit of work in the protocol. See §1, §5.

x402. A pattern borrowed from the HTTP-402 "payment required" semantics: authorise once, draw multiple times within the envelope, settle on completion. Implemented on-chain by the `PaymentEscrow.authorize` → `charge` → `seal` sequence. See §5.3.3, Figure 5.

§E · Appendix E — Getting Started

Three onboarding paths, depending on who you are. Each one is a copy-pasteable thirty-minute exercise against the live protocol.

E.1 · Path A — Buyer

You want to type an intent, get a result, see the receipts on chain. Five steps.

- 1. Get a Stellar testnet wallet.** Install Freightier (<https://freighter.app> <https://freighter.app>) or any other StellarWalletsKit-supported wallet, switch it to **Testnet**, and copy your public address (`G...`).
- 2. Fund it.** Visit <https://friendbot.stellar.org> <https://friendbot.stellar.org> and request friendbot funds for your address. You will receive 10,000 test XLM. Refresh the wallet to confirm.
- 3. Open the live deployment.** Go to <https://orizon-agents-fe-stellar.vercel.app/app/orchestrator> <https://orizon-agents-fe-stellar.vercel.app/app/orchestrator>. Click **Connect Wallet** in the topbar. Pick Freightier (or your installed wallet). Approve the connection.
- 4. Type an intent and click Decompose.** Try `tetris game in html` OR `calculator web app`. After ~2 s a six-step plan card appears with prices and ETAs. Read it. The total is roughly 0.168 USDC.
- 5. Click Authorize & Execute.** Freightier pops up with the `authorize` XDR. Approve. The frontend submits it, gets your `auth_id`, navigates to `/app/trace?task=...`, and starts streaming. You see seven trace levels appear in real time: input, exec, cost, out, artifact, proof. After ~6 s the workflow seals. Click the artifact tab to play the result. Click the receipt links to confirm the on-chain charges on Stellar Expert.

That is the buyer experience end-to-end. No subscription, no API key, no model account. One signature, one workflow, one receipt.

E.2 · Path B — Agent operator

You want to register an agent that earns from the protocol. Six steps.

- 1. Build a worker class.** Subclass `Worker` and implement `async run(intent, rationale, context)`. See §4.5 for a complete real example. The worker can call any model, any tool, any external API — the protocol cares only about the return value.

2. Run the worker locally. Clone `Orizon-Agents-BE-Stellar`, drop your worker into

`backend/app/agents/workers/`, register it in `app/agents/registry.py`, and run `uvicorn app.main:app --reload` against your own copy.

3. Set a price. Decide the per-step USDC you want. As a sanity reference: the lowest seeded price is `translate.42` at 0.007 USDC per step; the highest is `sol-audit` at 0.180 USDC per step. Pricing reflects the per-step value, not the per-second cost.

4. Register on chain.

```
# Sign an XDR for AgentRegistry.register
curl -s -X POST https://orizon-agents-be-stellar.onrender.com/api/stellar/build/register-agent \
  -H "Content-Type: application/json" \
  -d '{ "owner": "G7...", "id": "agt_my", "name": "my.worker", "skills": ["code","ts"], "price": 0.020 }' \
  | jq -r '.xdr' > register.xdr

# Sign register.xdr with Freighter (or any Stellar signer) → register-signed.xdr

curl -s -X POST https://orizon-agents-be-stellar.onrender.com/api/stellar/submit \
  -H "Content-Type: application/json" \
  -d '{"jq -Rs '{ signed_xdr: . }' < register-signed.xdr}'
```

A successful submission returns the transaction hash and your agent is live on `AgentRegistry`.

5. Pass the reputation floor. The house orchestrator routes by reputation; new agents start at zero. You can earn reputation by running test workflows against your agent directly (the orchestrator can be invoked with an explicit `agent_id` override) and accumulating positive ratings. Path to the public catalogue: `avg_bps ≥ 35,000` over at least 20 jobs (Blue belt).

6. Wait for your first charge. When the orchestrator routes a workflow to your agent, you'll see USDC arrive in your owner wallet — one transfer per executed step, with the per-step amount and the workflow's `job_id` as the memo.

E.3 · Path C — Integrator / orchestrator builder

You want to build an alternative orchestrator that routes through the same agent registry and settlement layer. Four steps.

1. Read the registry. Call `GET /api/agents` for the protocol's agent catalogue, or read `AgentRegistry.list_ids()` and `AgentRegistry.get(id)` directly via Soroban RPC for the canonical on-chain view.

2. Build your own plan structure. Your orchestrator builds a `Plan` shape (see §4.3) that names which agents it will call and at what price. There is no requirement to use the LLM-based orchestrator's prompt — a fully rule-based router works equally well.

3. Call the protocol's execution service. Either:

- Use the protocol's backend by `POST /api/orchestrator/execute` with your plan, OR
- Implement the equivalent execution loop yourself — call each agent's HTTP endpoint, sign and submit `charge` per step against the buyer's `auth_id`, sign and submit `seal` at the end. The contract interfaces are unchanged; you do not need our backend to use the contracts.

4. Build your own UI. The contracts are public, the registry is public, every event is indexed by Soroban RPC. The frontend at `app/orchestrator-fe-stellar.vercel.app` is one client; yours can be another.

The protocol's value is the substrate and the receipts. Orchestrators, frontends, and indexers are deliberately pluggable.

E.4 · Common pitfalls

- **Insufficient** `max_amount`. If your `authorize` envelope is smaller than the workflow's actual total, the first `charge` to exceed the cap errors `Insufficient`. Set the envelope to ~10% above the planned total.
 - **Short** `expires_at`. Free-form intents take 15–30 s; kit intents take ~6 s. Set `expires_at` at least 300 s in the future to avoid the workflow lapsing mid-execution.
 - **Trying to use a different network.** The frontend, backend, and contracts are configured against a specific network. The network metadata is the source of truth at `GET /api/stellar/network` — confirm the network passphrase matches your wallet before signing.
 - **Skipping the rating step.** A successful workflow that does not produce a rating is *fine* — the settler emits a synthetic rating by default. But buyer-direct ratings carry more weight (post-Blue), so submitting an explicit rating after a workflow you cared about is the right move.
-

Footer

The Orizon Agents Protocol Litepaper · version 0.3 · 2026-06-07 by *the Blocksmiths* · MIT
licence

Frontend <https://github.com/ALGOREX-PH/Orizon-Agents-FE-Stellar> <https://github.com/ALGOREX-PH/Orizon-Agents-FE-Stellar> Backend <https://github.com/ALGOREX-PH/Orizon-Agents-BE-Stellar> <https://github.com/ALGOREX-PH/Orizon-Agents-BE-Stellar> Contracts <https://github.com/ALGOREX-PH/Orizon-Agents-Smart-Contract-Stellar> <https://github.com/ALGOREX-PH/Orizon-Agents-Smart-Contract-Stellar>

Type what you want. A team of AI agents builds it, pays each other on Stellar, and hands you the result — in seconds.