

Orizon Agents — Week 2 Tranche Submission

Proof of Deliverables

Programme	Stellar Community Fund — Blue Belt Instawards (Cohort 2026)
Milestone	M2 · Week 2 — Reputation-Gated Routing (Deliverable D2) + External Agent Execution Path (Epic 2)
Sprint week	Mon 2026-09-14 → Fri 2026-09-18
Network	Stellar testnet only
Team	Danielle Bagaforo Meer — lead engineer (GitHub ALGOREX-PH) Rieselle Saure ("Rie") — PM + QA (GitHub rie-hash14)
Evidence captured	2026-09-19, from the live testnet deployment and public GitHub / Stellar Expert pages. Each page states its source URL and capture time.

DEPLOYED TESTNET CONTRACT IDS

AgentRegistry	CAPHXWU53UZUZJGV7IAE57NNMH3YYB5MTW06YA53KKMXSFVL0ITBJ3GQ
ReputationLedger	CDCS0BEVZUPQZV5GV4D6KYHZCLNGW2KXY74RUHSZ3EZUXF34DPW422ZT
PaymentEscrow	CBJPTMAPMG0DGZCZ2IMEQSRUX3WGUXNMKDTNN2KMJ3NF6YZ50J5525PI
AttestationRegistry	CBYUZXK0ET43UXTBXZUJIBBJW50DGD2J2AZVVXCR3Q0NG0CAH0XQQHEGK
Asset SAC	CDLZFC3SYJYDZT7K67VZ75HPJVIEUVNIXF47ZG2FB2RMQQVU2HH6CYSC

EVIDENCE IN THIS DOCUMENT

–	What we shipped in Week 2	summary	p. 2
–	Stellar integration this week	summary	p. 3
01	Plan card: each step's reputation and the routing floor, before payment	3.04 · D2	p. 4
02	Marketplace standing: floor stated once, provenance and standing marks	3.05 · D2	p. 5
03	Operator dashboard ("My Agents"), no wallet connected	2.06 · Epic 2	p. 6
04	Bind an endpoint to an on-chain agent, by wallet signature	2.01 / 2.05 · Epic 2	p. 7
05	Backend /readiness: cold-start headroom and whether ratings can be written	BE #58 / #59 · D2	p. 8
06	Backend PR #59 — planner-outage fallback, observable ratings, log redaction	BLO-121 · backend	p. 9
07	Frontend PR #62 — the plan card says when a plan is the planner fallback	BLO-121 · frontend	p. 10
08	UAT PR #3 — Rie's Week-2 QA commits	QA · rie-hash14	p. 11
09	Backend PRs merged 13–18 Sep 2026 (13)	backend delivery	p. 12
10	Frontend PRs merged 13–18 Sep 2026 (10)	frontend delivery	p. 13
11	Contracts PR #2 — the deployed address book is tracked in git	3.07	p. 14
12	The public reference external agent	2.04 · Epic 2	p. 15
13	set_scorer on the ReputationLedger — on-chain rating writes enabled for production	on-chain ratings enabled	p. 16
14	Week-2 public build post on X	7.02 · published 2026-09-19	p. 17

Every page names the public URL it was captured from, so each claim can be checked live. On the plan-card and marketplace pages every agent carries the ≈3.50 starting estimate and clears the 2.75 routing floor, which each page states and applies.

What we shipped in Week 2

This week we shipped the **external agent execution path** and **reputation-gated routing (Deliverable D2)**, both live on Stellar testnet at orizons.xyz.

External agent execution path — Epic 2

- **Endpoint binding by wallet signature** (2.01, 2.05) — an operator attaches their service URL to their on-chain agent by signing a one-time challenge with the wallet that owns it, on its own or inside registration.
- **Signed, bounded dispatch** (2.02) — every workflow step sent to an external agent is signed so the operator can verify it came from Orizon; unsafe URLs are refused and every step has a deadline and a size cap.
- **Fair failure rules** (2.03) — a step that times out or returns unusable output is not charged and counts against the agent's reputation.
- **Reference agent** (2.04) — a public, copyable example agent that runs from its README.
- **Operator dashboard** (2.06) — each operator's agents, earnings and reputation in one place, with transaction links.

Reputation-gated routing — Epic 3 · Deliverable D2

- **A reputation floor on every plan** (3.01, 3.02) — every planning path reads each agent's on-chain reputation and applies the floor on its conservative lower bound.
- **The trust signal before payment** (3.04) — the plan card shows each agent's reputation and its source, the floor applied, and any agent the floor acted on, with the reason.
- **Honest reads** (3.03) — if reputation cannot be read from chain, the plan and the card say so.
- **Marketplace standing** (3.05) — the floor stated once, on-chain agents marked, standing marks per agent.
- **New agents are hireable on day one** (3.06) — a brand-new agent clears the floor by 177 bps, guaranteed by a test and a startup check.
- **Contract addresses kept in agreement** (3.07) — checked daily and against the live deployment.
- **On-chain ratings from production** — every paid run is rated, and `/readiness` verifies production is the ledger's authorized scorer.

Fixes and hardening

- A planning-service outage now returns a fallback plan instead of an error (BLO-121).
- The floor holds after planning: a step survives only if its agent was offered to the planner.
- Secrets are masked in every log line; prices show the network's real currency (XLM); accessibility and mobile fixes across the app.

QA — Rie

- End-to-end QA on testnet of the external execution path (6.05) and the operator surfaces (6.06), a reputation-floor test suite, and 14 defects logged and handed to engineering.

27

pull requests merged

5

public repositories

1,217authored commits (Dan
1,080 · Rie 137)**1,534**

backend tests passing

1,086

frontend tests passing

Stellar integration this week

All work runs on **Stellar testnet**.

SDKs and libraries

- **Backend** — Python `stellar-sdk` 13.2.1: builds unsigned XDR for wallets to sign, simulates, prepares and submits Soroban transactions, reads contract storage, verifies SEP-53 message signatures.
- **Frontend** — `@stellar/stellar-sdk` 15.0.1 and `@creit.tech/stellar-wallets-kit` 2.1.0 for wallet connection and transaction / message signing.
- **Reference agent** — verifies Orizon's SEP-53 dispatch signatures with PyNaCl (ed25519) and hand-rolled StrKey decoding, so operators need no Stellar SDK.

Soroban contracts

- **AgentRegistry** — `owner_of` gates endpoint binding; `set_active` delisting is honoured by routing; `register` for new agents.
- **ReputationLedger** — batched, time-bounded `rep_state` reads feed the routing floor; `submit` records ratings; `/readiness` reads its storage via `getLedgerEntries` to confirm the authorized scorer.
- **PaymentEscrow** — the buyer authorizes a workflow cap from their own wallet; the operator dashboard reads the escrow's on-chain events.
- **AttestationRegistry** — sealer role assigned to the production key. **Native XLM SAC** — the escrow asset.

Wallet integration

- **Binding by signature** — the owner wallet signs a SEP-53 challenge naming the agent, endpoint and a nonce; a separately signed message revokes it.
- **Signed dispatch** — every step is signed over SEP-53 by a dedicated keypair published as `dispatch_signer` (`GB5MKHDF...KCMR`).

Transactions (testnet)

Date	Call	Transaction hash
09-17	AgentRegistry.set_active (delist)	a710b6776042d810fa1a41ec17cc0b299c606fc2a7f54bd28c129e20bbe6e8e4
09-17	AgentRegistry.register (calculatorai)	0741a0822b6976f88a4582ffc65f1528004a9a5c3c544171e4be7ba099b1c8aa
09-17	PaymentEscrow.authorize (0.168 XLM)	67701b46ef60bf488481464cf527aeadb01bb407cf294efc120c5f4ecc31e56d
09-17	AgentRegistry.register (algorex)	7e3b6c02731906ddb685b080f11c63e0c1c2665a7bad0f0836bfa7ee5d83c872
09-19	ReputationLedger.set_scorer	216e1b5f6ade4d75ec671bcda27b462bfd373d041b1ba2150d76002ee8d201f8
09-19	AttestationRegistry.set_sealer	c965980fd06d5917bfa46dfc72898422a3f50136e0ac4f487e4ed0f7a19a3c

Each hash opens at `stellar.expert/explorer/testnet/tx/<hash>`.

APIs

- **Soroban RPC** (`soroban-testnet.stellar.org`) — `simulateTransaction`, `sendTransaction`, `getLedgerEntries`, `getEvents`. **Horizon** testnet for transaction verification.
- The contract address book is tracked in git; a daily check and a 6-hourly production smoke test compare the live `/api/stellar/network` ids with it.

Plan card: each step's reputation and the routing floor, before payment

The live orchestrator decomposed the demo intent **"tetris game in html"** (chosen with the page's own preset button, then *Decompose*) into a six-step plan. Before the buyer is asked to authorize anything, the card states the routing floor that was applied (**floor 2.75 · applied**, checked against each agent's reputation lower bound) and shows every step's agent with its reputation and that reputation's source: the **≈★ 3.50** chip, where **≈** marks a prior estimate rather than settled on-chain ratings.

Agents on the live registry carry the **≈3.50 starting estimate**, which clears the 2.75 floor, so the floor summary reads "the floor acted on no agents". The five agents it lists as having no endpoint bound are not candidates because there is nothing to dispatch a step to. Nothing was authorized, simulated or paid for this capture.

Source <https://orizons.xyz/app/orchestrator>

Captured 2026-09-19 15:47 PHT · screenshots/01-orizons-plan-card.png

The screenshot displays the Orizon Agents Orchestrator interface. The left sidebar contains a navigation menu with options: Overview, Agents, Register, Bind, My Agents, Reputation, Orchestrator (selected), Trace, Events, Send XLM, PDAX Ramp, Wallet, and Flow. A 'NETWORK' status box shows 2,498 agents online with 94% completion. The main area is titled 'Orchestrator' and shows the intent 'tetris game in html'. Below the intent, there are buttons for 'TETRIS GAME IN HTML', 'CALCULATOR WEB APP', 'SNAKE GAME IN HTML', and 'POMODORO TIMER WITH SOUND', along with a 'DECOMPOSE' button. The 'Execution plan' section shows a plan with 6 steps. The first step is 'ROUTING FLOOR' with a status of 'FLOOR 2.75 · APPLIED'. The plan details state: 'Every agent considered for this plan was checked against a 2.75 routing floor before the planner chose. The check is each agent's reputation lower bound against that floor, never its headline score. 6 steps planned · the floor acted on no agents · 5 agents with no endpoint bound were never candidates'. The plan steps are listed in a table:

Step	Agent	Reputation	Action	Estimate	ETA
01	RESEARCH.PRO	≈★ 3.50	extract feature brief + edge cases for the build	0.024	0.6s
02	SEO.BRIEF	≈★ 3.50	produce brand identity: name, tagline, audience	0.009	0.5s
03	DESIGN.FIGMA	≈★ 3.50	lock design tokens: palette, typography, motion	0.038	0.4s
04	CODE.GEN	≈★ 3.50	implement single-file HTML using brief + tokens	0.054	2.6s
05	CODE.CRITIC	≈★ 3.50	polish pass: aily, motion, persistence, edge cases	0.052	1.8s
06	DEPLOY.VO	≈★ 3.50	seal artifact + record on-chain proof	0.011	0.4s

Below the steps, a 'REPUTATION FLOOR · NO CHANGES' section states '5 with no endpoint bound'. At the bottom, a 'WALLET REQUIRED' section prompts the user to 'Connect Freightier (testnet) to pay with x402 on-chain, or run a simulated pass.' and includes buttons for 'CONNECT WALLET', 'PAY WITH FIAT', and 'SIMULATE'.

Marketplace standing: floor stated once, provenance and standing marks

The Agent Registry states the selection floor **once**, above the table (**floor 2.75**, checked against each agent's reputation lower bound, never the headline score). Agents registered on-chain against the public registry carry an **"external"** provenance mark; the seeded first-party catalog rows carry none. Standing marks sit beside the agent: **"not yet operational"** (registered, no endpoint bound) and **"delisted by operator"**.

A third mark, **"below floor · not eligible"**, flags any agent whose lower bound falls under the floor; every agent on the live registry currently clears it.

Source <https://orizons.xyz/app/agents>

Captured 2026-09-19 15:47 PHT · screenshots/02-orizons-agents-standing.png · Shown in excerpts (slices of the full-page capture, 1440 × 2057 px); nothing inside a slice is altered.

ORIZON AGENTS

CONSOLE / AGENTS

CONNECT WALLET LIVE

WORKSPACE

- Overview
- Agents
- Register
- Bind
- My Agents
- Reputation
- Orchestrator
- Trace
- Events
- Send XLM
- PDAX Ramp
- Wallet
- Flow

Agent Registry

ERC-8004 profiles — identity, skills, price, reputation.

+ REGISTER AGENT

Selection floor FLOOR 2.75

The floor decides which agents the orchestrator will consider when it builds a plan. It is checked against each agent's reputation lower bound — what a score is worth once it has been discounted for how few rated jobs stand behind it — and never against the headline score in the reputation column. An agent can show a strong score and still sit below the floor.

search name or skill...

ALL ROUTABLE ONLINE IDLE OFFLINE

ID	AGENT	SKILLS	PRICE / REPUTATION	CALL	RUNS	STATUS	
agt_01h8	copywrite.v3	COPY SEO EN	0.012	3.50	18,420	ONLINE	VIEW
agt_02k2	design.figma	UI TOKENS FIGMA	0.018	3.50	7,321	ONLINE	VIEW

... seeded catalog rows agt_03d9 → agt_12r0 omitted here (no marks) — all rows are in the full-page PNG ...

w1_audit_a7x	week1 audit scratch (non-executing)	AUDIT SCRATCH	0.001	3.50	0	ONLINE	VIEW
sign_probe_bb5c121.05	sign probe	PROBE SIGN	0.021	3.50	0	ONLINE	VIEW
Testing_Agent	Testing Agent		0.030	3.50	0	ONLINE	VIEW
dan_w1_probe	Dan Week 1 Registration	DEMO	0.050	3.50	0	OFFLINE	VIEW
spike_97437	Spike Agent	SPIKE	0.010	3.50	0	ONLINE	VIEW
uat605_ext_op	UAT 6.05 haiku operator	HAIKU POETRY UAT605	0.010	3.50	0	ONLINE	VIEW
calculatorai	Calculator AI	ADDING SUBTRACTING	0.010	3.50	0	ONLINE	VIEW
algorex	algorex		0.010	3.50	0	ONLINE	VIEW


Operator dashboard ("My Agents"), no wallet connected

The operator dashboard as a visitor sees it with **no wallet connected**. Agent ownership is read from the chain, so the page shows only a prompt to connect; connecting reads public registry data and signs nothing.

Once the owner connects their wallet, the same page lists every agent that wallet owns on-chain, with its reputation and routing standing, its endpoint binding, and a *Settlement* panel of what it has been paid (read from the escrow's on-chain charge events).

Source <https://orizons.xyz/app/operator>

Captured 2026-09-19 15:47 PHT · screenshots/03-orizons-operator-dashboard.png



Bind an endpoint to an on-chain agent, by wallet signature

The operator endpoint-binding page. An operator enters the id of the agent they registered on-chain and the HTTPS endpoint that runs its work; the registry issues a one-time challenge naming the agent, the endpoint and a nonce, and the **wallet that owns the agent signs it** — nothing else can authorize the binding. Ownership stays on-chain; the endpoint is an off-chain service record that can be re-bound.

Captured with no wallet connected, so the button reads “connect the owner wallet to sign”. `weather_bot` and `https://agent.example.com/run` are the form's placeholders, not a real binding.

Source <https://orizons.xyz/app/bind>

Captured 2026-09-19 15:47 PHT · screenshots/04-orizons-bind.png

ORIZON AGENTS

CONSOLE / BIND

CONNECT WALLET LIVE

Bind an Endpoint

Point one of your agents at the HTTPS endpoint that runs its work. The wallet that owns the agent signs the binding — nothing else can authorize it.

CONNECT WALLET

ENDPOINT BINDING

SIGNING WALLET · CONNECT A WALLET

AGENT ID

weather_bot

the id you registered on-chain · letters, digits and underscore

CURRENT BINDING

enter an agent id to see what it points at today

ENDPOINT URL

https://agent.example.com/run

where the agent receives work · https is assumed when you omit the scheme

TRUST BOUNDARY

Ownership of this agent is on-chain and permanent. The endpoint is not: it is an off-chain service record held by the registry, so you can move hosts by binding again, and nothing about your ownership changes.

BIND ENDPOINT connect the owner wallet to sign

HOW BINDING WORKS

1. The registry issues a one-time challenge naming your agent, the endpoint and a nonce.
2. Your wallet signs that exact string. The signature proves the binding came from the account that owns the agent — no password, no account to create.
3. The challenge is short-lived. If it expires while the wallet popup is open, a fresh one is requested rather than a dead signature submitted.

No agent yet? [Register one first](#)

WORKSPACE

- Overview
- Agents
- Register
- Bind
- My Agents
- Reputation
- Orchestrator
- Trace
- Events
- Send XLM
- PDAX Ramp
- Wallet
- Flow

NETWORK

2,498 agents online
avg completion 94%

operator
testnet

Backend /readiness: cold-start headroom and whether ratings can be written

Captured 2026-09-19

The production backend's readiness probe. `cold_start` shows that a brand-new agent clears the routing floor: its reputation lower bound is **5677 bps against a 5500 bps floor — a 177 bps margin** (2.84 vs 2.75 on the dApp's 5-point scale), so `routable: true`. `ratings.writer` reports whether production can write ratings; it reads `scorer`: the backend's signer is the ReputationLedger's authorized scorer (`GDB4N25...`), so every paid run's ratings are written on-chain.

The JSON body was re-indented in the browser for legibility; its content is unchanged.

Source <https://orizon-agents-be-stellar.onrender.com/readiness>

Captured 2026-09-19 15:47 PHT · screenshots/05-be-readiness.png · Shown: the top 520 px of the 1440 × 900 frame; the rest of it is blank.

GET https://orizon-agents-be-stellar.onrender.com/readiness — fetched 2026-09-19T06:59:37.231Z (response body, indented for legibility)

```
{
  "status": "ready",
  "llm": "ok",
  "stellar": "configured",
  "signer": "configured",
  "pdax": "configured",
  "cold_start": {
    "routable": true,
    "lower_bound_bps": 5677,
    "floor_bps": 5500,
    "margin_bps": 177
  },
  "ratings": {
    "writer": "scorer",
    "signer": "GDB4N25UYM3YNTTAWX7LSGI2P70R62QZXRQWAGF5TFVENDKCTTCDHP",
    "scorer": "GDB4N25UYM3YNTTAWX7LSGI2P70R62QZXRQWAGF5TFVENDKCTTCDHP"
  }
}
```

Backend PR #59 — planner-outage fallback, observable ratings, log redaction

The backend half of the Week-2 branch `feat/week-2-dan` (124 commits, merged 2026-09-18). **BLO-121**: when the planner (LLM) fails, the buyer now gets a fallback plan built from agents that passed the routing checks, flagged `planner_fallback: true`, instead of a 502. **Observable ratings**: `/readiness` gains the `ratings` field, the service logs at boot whether it can write ratings, and a failed rating names its cause. **Log redaction**: configured secrets and key-shaped strings are masked on every log line.

Source <https://github.com/BIOcksmiths/Orizon-Agents-BE-Stellar/pull/59>
Captured 2026-09-19 15:47 PHT · screenshots/06-be-pr-59.png · Shown in excerpts (slices of the full-page capture, 1440 × 4818 px); nothing inside a slice is altered.

The screenshot shows a GitHub pull request page for the repository `BIOcksmiths/Orizon-Agents-BE-Stellar`. The pull request is titled "Week 2: planner outage fallback (BLO-121), observable ratings, log redaction #59" and is in the "Merged" state. It was merged by `ALGOREX-PH` 19 hours ago, merging 124 commits from the `feat/week-2-dan` branch into the `main` branch. The page shows a comment from `ALGOREX-PH` detailing the changes and fixes for BLO-121. The comment includes a list of bullet points describing the bug, the planner's answer, timeouts, new response fields, provider text, and ratings observability. It also mentions that secrets are now redacted and that two previously misreported outcomes are fixed. The comment concludes with a note about test-hygiene and verification notes being omitted in the full-page PNG. The right sidebar shows the "Reviewers" section with "No reviews" and the "Assignees" section with "No one assigned". The "Labels" section shows "None yet". The "Projects" section shows "None yet". The "Milestone" section shows "No milestone". The "Development" section shows "Successfully merging this pull request may close these issues. None yet". The "1 participant" section shows the profile of `ALGOREX-PH`. The bottom of the page shows the merge commit `5822729` into `main` 19 hours ago, with a "View details" button.

Platform Solutions Resources Open Source Enterprise Pricing Search Sign in Sign up

BIOcksmiths / Orizon-Agents-BE-Stellar (Public) Notifications Fork 0 Star 1

<> Code Issues Pull requests 9 Actions Projects Security and quality Insights

Week 2: planner outage fallback (BLO-121), observable ratings, log redaction #59

Merged ALGOREX-PH merged 124 commits into main from feat/week-2-dan 19 hours ago

Conversation 0 Commits 124 Checks 1 Files changed 27 +2,875 -47

ALGOREX-PH commented 19 hours ago • edited Collaborator

Week 2 work for Dan: fixes **BLO-121**, and makes ratings observable, so a deployment that can't write ratings says so instead of failing silently.

BLO-121: a planner outage now gets a modest plan, not a 502

- The bug**: when the LLM call failed, agno put the provider's error `string` into `result.content`. `plan: Plan = result.content` then crashed, and the router returned **502 decompose_failed**.
- Checking the planner's answer**. `planner_plan()` only accepts a real `Plan` from a run agno didn't mark errored or cancelled.
 - An error string, a cancelled run, an unparsed answer, or a raised provider exception all degrade to the existing fallback plan.
 - That fallback comes from the agents offered to the planner, with `degraded` set by the floor rule.
- Timeouts are unchanged**: they still return **504 decompose_timeout**. Nothing offered is still **503 no_routable_agents**.
- New `DecomposeResponse.planner_fallback`** (additive, default false): true whenever the served plan is the fallback rather than the planner's own. The FE PR shows it to the buyer.
- The provider's text is never returned**. It is logged at **WARNING** with secrets redacted.
- Verified end to end with the ticket's own repro (blank key, "write a haiku about the sea"): **200**, `planner_fallback: true`.

Ratings are observable

The live ledger holds zero ratings, and nothing in the service could say why. The diagnosis is on BLO-41. Now:

- `/readiness` gains `ratings`**: `{ "writer": "disabled|no_signer|scorer|not_scorer|unchecked", "signer": "...", "scorer": ... }`.
 - `scorer` is read from the ledger's instance storage and cached for 300 s, with failed reads retried after 30 s.
 - The probe itself makes no network call.
 - It is informational and never changes readiness.
 - Checked read-only against testnet: the local signer **is** the stored scorer.
- Startup line**: once at boot, the service logs whether it can write ratings. Anything other than `scorer` is logged at **WARNING**.
- No more silent skip**. `submit_ratings` used to `return` without a word when unconfigured. It now logs a rate-limited **WARNING** and traces `ratings not submitted: <reason>` on the task.
- Failed ratings name their cause**. A failed rating now reads `reputation submit failed for X: Unauthorized | NotFound | Replay | OutOfRange | contract error #N | rpc error`, and never echoes exception text.
- Two previously misreported outcomes fixed**: a rating whose transaction came back **FAILED** or timed out used to be traced as "rated N/100". It is now an error.

Secrets never reach a log line

- Masking**: `SecretRedactionLogFilter` on the root handler masks every configured secret by exact value, longest first. It also masks anything shaped like an API key (including provider-starred keys) or a Stellar secret seed, in messages, args and tracebacks. Public keys, contract ids and transaction hashes are untouched.
- agno's logs**: agno installs its own Rich handler with propagation off, so its lines bypassed the JSON format, the request id and any redaction. That includes the verbatim provider error it logs at **ERROR**. Its loggers are now routed through the root handler, at **WARNING**.
- One rule for the planner**: `loggable` now uses the same shared redaction.

... test-hygiene and verification notes and the 124-commit timeline omitted — all in the full-page PNG ...

ALGOREX-PH merged commit 5822729 into main 19 hours ago 1 check passed View details

Frontend PR #62 — the plan card says when a plan is the planner fallback

The frontend half of BLO-121 (34 commits, merged 2026-09-18). When the backend serves a fallback plan, the plan card shows a calm notice above Authorize and Simulate saying the planner was unavailable and the plan was built only from agents that passed the routing checks, with **"Ask the planner again"**, which re-runs the intent the buyer actually submitted. The description lists its own verification (unit, Playwright and axe checks); the merge commit shows 5 checks passed.

Source <https://github.com/BIOcksmiths/Orizon-Agents-FE-Stellar/pull/62>
Captured 2026-09-19 15:47 PHT · screenshots/07-fe-pr-62.png · Shown in excerpts (slices of the full-page capture, 1440 × 3317 px); nothing inside a slice is altered.

The screenshot shows a GitHub pull request interface. At the top, the repository is 'BIOcksmiths / Orizon-Agents-FE-Stellar'. The pull request title is 'Week 2: tell the buyer when a plan is the planner fallback (BLO-121) #62'. It is marked as 'Merged' and shows 'ALGOREX-PH merged 34 commits into main from feat/week-2-dan' 19 hours ago. The pull request has 34 commits, 4 checks, and 10 files changed. The description of the pull request is as follows:

The frontend half of BLO-121. When the backend serves a **fallback plan** because the planner failed or answered with nothing usable (`planner_fallback: true` , added in BE [#59](#)), the buyer is now told so before paying.

What changed

- **Types and guard.** `DecomposeResponse.planner_fallback?: boolean` . The guard accepts true, false, null or absent, and rejects anything else, in the same way it handles `reputation_degraded` .
- **PlannerFallbackNotice on the plan card**, above Authorize and Simulate:
 - A calm, informational notice (`role="status"` , never an alert). It uses violet, because magenta means a warning on this card, and a glyph plus text so the meaning never rests on colour alone.
 - It says the planner was unavailable, so this is a minimal fallback plan built only from agents that passed the routing checks.
 - It never names a cause and never discourages paying.
- **"Ask the planner again"** re-runs the page's existing decompose with the intent the buyer actually submitted, not whatever the intent box says by then. No new API. It is disabled while a simulate or authorize run is in progress.
- **Accessibility.** Authorize's `aria-describedby` lists every notice on the page, in reading order: this one, then the unverified-reputation banner. When neither is shown, the attribute is left off entirely.

Tests

- **Unit:** 29 new tests (notice, plan card, guard).
- **E2E:** new `e2e/plan-fallback.spec.ts` :
 - the notice appears above the actions, and doesn't appear when the flag is false or absent;
 - the describedby combinations;
 - axe WCAG 2.0/2.1 A+AA is clean in both states;
 - at 390px there's no sideways scroll and nothing is clipped;
 - the retry asks the same intent again, and the planner's plan replaces the fallback.

Verification (on the merged branch)

- `typecheck` , `lint` , and `Prettier` on tracked files: clean.
- `vitest --coverage` : **981/981**. Coverage 96.23 / 91.88 / 94.63 / 98.44 on `lib/**` .
- **Playwright:** **105 passed, 0 failed** (from the summary line; no **X**).
- `build` , `bundle` (every route within 480 KB), `check:addresses` , `test:parity` 20/20, and `lighthouse` : pass.

... the 34-commit timeline and the Vercel preview comment omitted — both in the full-page PNG ...

ALGOREX-PH merged commit `3070f5c` into `main` 19 hours ago
5 checks passed

UAT PR #3 — Rie's Week-2 QA commits

Rie's Week-2 QA work in the public UAT repository, merged as PR #3 on 2026-09-17: **142 commits, 141 of them by rie-hash14** (21 on 12 Sep, 50 on 16 Sep, 70 on 17 Sep) plus one merge of `main` by ALGOREX-PH. They add the UAT coverage for **6.02** (reputation floor), **6.05** (external dispatch) and **6.06** (operator surfaces), and the defect log entries **D-036 → D-049**.

Source <https://github.com/BI0cksmiths/Orizon-Agents-UAT-Stellar/pull/3/commits>

Captured 2026-09-19 15:47 PHT · screenshots/08-uat-pr-3-rie-commits.png · Shown in excerpts (slices of the full-page capture, 1440 × 9837 px); nothing inside a slice is altered.

The screenshot displays the GitHub interface for a pull request titled "Uat #3". The repository is "BI0cksmiths / Orizon-Agents-UAT-Stellar". The pull request is merged and shows 142 commits. The commits are grouped by date: Sep 12, 2026; Sep 16, 2026; and Sep 17, 2026. Each commit entry includes a description, the author (rie-hash14), the commit hash, and a link to the commit. The pull request is authored by ALGOREX-PH and merged into the 'main' branch from the 'uat' branch.

Uat #3

Merged ALGOREX-PH merged 142 commits into `main` from `uat` 2 days ago

Conversation 0 Commits 142 Checks 0 Files changed 25 +4,217 -6

Commits on Sep 12, 2026

- updated uat test plan module (f75d289) rie-hash14 committed last week
- updated uat test plan module (c75823e) rie-hash14 committed last week
- updated uat test plan module (bda5a74) rie-hash14 committed last week
- updated uat test plan module (57ca660) rie-hash14 committed last week
- added uat lockfile module (b3df188) rie-hash14 committed last week
- updated uat traceability module (f68bf37) rie-hash14 committed last week

... remaining 15 commits of 12 Sep omitted ...

Commits on Sep 16, 2026

- updated reputation floor spec module (a461492) rie-hash14 committed 3 days ago
- updated uat defects module (2849ce5) rie-hash14 committed 3 days ago
- updated uat defects module (3bf60a7) rie-hash14 committed 3 days ago

... remaining 47 commits of 16 Sep omitted ...

Commits on Sep 17, 2026

- updated tsconfig module (903d916) rie-hash14 committed 2 days ago
- added stellar key module (35aeb8b) rie-hash14 committed 2 days ago
- added verify dispatch module (3ef88de) rie-hash14 committed 2 days ago

... the 17 Sep commits in between omitted — every commit is in the full-page PNG ...

- merged operator surfaces spec module (700abc9) rie-hash14 committed 2 days ago
- Merge branch 'main' into uat (a46523c) ALGOREX-PH authored 2 days ago (Verified)

Backend PRs merged 13–18 Sep 2026 (13)

GitHub's own search for backend PRs merged 13–18 Sep 2026 returns **13**, numbered #41 to #59, covering Epic 2 (2.01 endpoint binding, 2.02 verifiable bounded dispatch, 2.03 external step failure semantics, 2.04 reference-agent contract, 2.06 settlement evidence), Epic 3 (3.02 excluded agents and the floor, 3.03 routing-time reads) and the Week-2 branch.

Source	https://github.com/BI0cksmiths/Orizon-Agents-BE-Stellar/pulls?q=is%3Apr+is%3Amerged+merged%3A2026-09-13..2026-09-18
Captured	2026-09-19 15:47 PHT · screenshots/09-be-pull-requests.png

PlatformSolutionsResourcesOpen SourceEnterprisePricing

Search / Sign in Sign up

BI0cksmiths / Orizon-Agents-BE-StellarPublic

Notifications Fork 0 Star 1

<> Code Issues

Pull requests 9

Actions

Projects

Security and quality

Insights

Q is:pr is:merged merged:2026-09-13..2026-09-18

Labels 12 Milestones 0

New pull request

Clear current search query, filters, and sorts

13 Total

AuthorLabelProjectsMilestonesReviewsAssigneeSort

Week 2: planner outage fallback (BLO-121), observable ratings, log redaction

#59 by ALGOREX-PH Collaborator was merged 19 hours ago

Epic 3 hardening: hold the floor after planning, enrich plan steps, ops fixes

#58 by ALGOREX-PH Collaborator was merged yesterday

Epic 2 hardening — kill switch, dispatch security, rating semantics

#57 by ALGOREX-PH Collaborator was merged 3 days ago

Week 2 — Epic 3 backend (consolidates #50, #53, #54, #55)

#56 by ALGOREX-PH Collaborator was merged 3 days ago

3.03 — Routing-time read hardening and degradation visibility

#52 by ALGOREX-PH Collaborator was merged 3 days ago

3.02 — Surface excluded agents and the floor

#51 by ALGOREX-PH Collaborator was merged 3 days ago

2.06 — Settlement evidence endpoint

#49 by ALGOREX-PH Collaborator was merged 3 days ago

Story 2.04: reference agent contract, published deadline, and the settlement finding

#48 by ALGOREX-PH Collaborator was merged 3 days ago

Story 2.03: external step failure semantics and rating consequences

#47 by ALGOREX-PH Collaborator was merged 4 days ago

Story 2.02: verifiable, bounded external dispatch

#45 by ALGOREX-PH Collaborator was merged 4 days ago

Harden the external step path: crash, rating and planner-injection fixes

#44 by ALGOREX-PH Collaborator was merged 4 days ago

Story 2.01: bind an operator endpoint to an on-chain agent id

#43 by ALGOREX-PH Collaborator was merged 4 days ago

Backend production hardening: pin money-route auth + guard consistency

#41 by ALGOREX-PH Collaborator was merged 4 days ago

ProTip! Mix and match filters to narrow down what you're looking for.

© 2026 GitHub, Inc.

TermsPrivacySecurityStatusCommunityDocsContactManage cookiesDo not share my personal information

Frontend PRs merged 13–18 Sep 2026 (10)

The same search on the frontend returns **10**, numbered #48 to #62, including the 2.01 binding UI (#50), the 2.05 binding step in registration (#51), the 2.06 operator dashboard (#52), the 3.02 floor visibility contract (#54), the 3.04 plan card (#56) and the Week-2 branch (#62).

Source	https://github.com/BIOcksmiths/Orizon-Agents-FE-Stellar/pulls?q=is%3Apr+is%3Amerged+merged%3A2026-09-13..2026-09-18
Captured	2026-09-19 15:47 PHT · screenshots/10-fe-pull-requests.png

PlatformSolutionsResourcesOpen SourceEnterprisePricing

Search /

Sign in

Sign up

BIOcksmiths/Orizon-Agents-FE-StellarPublic

NotificationsFork0Star1

<> CodeIssuesPull requests8ActionsProjectsSecurity and qualityInsights

Q is:pr is:merged merged:2026-09-13..2026-09-18

Labels12Milestones0

New pull request

Clear current search query, filters, and sorts

10 Total

AuthorLabelProjectsMilestonesReviewsAssigneeSort

Week 2: tell the buyer when a plan is the planner fallback (BLO-121)

#62 by ALGOREX-PH Collaborator was merged 19 hours ago

Epic 3 hardening: green CI, honest reputation surfaces, live contract parity

#61 by ALGOREX-PH Collaborator was merged yesterday

Epic 2 hardening — rebind coverage and transaction links

#60 by ALGOREX-PH Collaborator was merged 3 days ago

Week 2 — Epic 3 frontend (consolidates #53, #55, #57, #58)

#59 by ALGOREX-PH Collaborator was merged 3 days ago

3.04 — Plan card: reputation, source and exclusions

#56 by ALGOREX-PH Collaborator was merged 3 days ago

3.02 — Floor visibility contract

#54 by ALGOREX-PH Collaborator was merged 3 days ago

2.06 — Operator dashboard

#52 by ALGOREX-PH Collaborator was merged 3 days ago

Story 2.05: endpoint binding step in the registration flow

#51 by ALGOREX-PH Collaborator was merged 3 days ago

Story 2.01: operator endpoint binding UI

#50 by ALGOREX-PH Collaborator was merged 4 days ago

Frontend production hardening: a11y, skills-input UX, SSE guard

#48 by ALGOREX-PH Collaborator was merged 4 days ago

ProTip! What's not been updated in a month: [updated:<2026-08-19](#).

© 2026 GitHub, Inc. TermsPrivacySecurityStatusCommunityDocsContactManage cookiesDo not share my personal information

Contracts PR #2 — the deployed address book is tracked in git

Merged 2026-09-16. The contract repo's record of what is deployed (`addresses.json` for testnet, `addresses.mainnet.json`) was git-ignored, so it existed only on the machine that last deployed. It is now **tracked in git**, giving story 3.07's contract-id drift check a canonical file to compare the backend configuration against. The PR records that all five contract ids on both networks matched the backend's `.env.example` when committed.

Source <https://github.com/BIOcksmiths/Orizon-Agents-Smart-Contract-Stellar/pull/2>
Captured 2026-09-19 15:47 PHT · screenshots/11-contracts-pr-2.png

The screenshot shows a GitHub pull request (PR) titled "Track the deployed address book #2" in the repository "BIOcksmiths/Orizon-Agents-Smart-Contract-Stellar". The PR is merged and was created 3 days ago. It shows 3 commits and 3 files changed. The comment from ALGOREX-PH explains the purpose of the PR: to track the deployed address book in git, ensuring that the contract IDs are consistent across different environments (testnet, mainnet) and preventing drift. The comment includes a list of tracked files and a diff showing the changes to the address book files. The PR is referenced by several other PRs, including "3.07 — Catch contract address drift" and "Week 2 — Epic 3 backend (consolidates #50, #53, #54, #55)".

Track the deployed address book #2

ALGOREX-PH merged 3 commits into `main` from `feat/3.07-track-addresses` 3 days ago

Conversation 0 Commits 3 Checks 0 Files changed 3 +38 -2

ALGOREX-PH commented 3 days ago

Story 3.07 (BLO-111) asks for a check that catches contract-id drift between this repo and the backend config. That check could not exist, because **the file it would check against was not in this repo.**

```
38 tracked files · addresses tracked: 0
.gitignore:7 addresses.json
.gitignore:8 addresses.*.json
```

`addresses.json` and `addresses.mainnet.json` are this repo's record of which contracts are deployed, and ignoring them meant that record existed only on whichever laptop last ran a deploy. A fresh clone had no address book at all. No CI anywhere could compare against one. And the authoritative copy was one disk failure away from being recoverable only by reading it back out of a *downstream* copy in the backend's `.env.example`.

The ticket says the drift is "invisible to review and to CI by construction". This is the construction.

Nothing in these files is secret. Contract ids are public on-chain identifiers; the `admin` field is a public G-address. Every value is already committed to `README.md` here, and to `.env.example`, `README.md`, `onchain-details.tsx` and `lib/api.test.ts` in the other two repos — verified before committing, along with a check that neither file contains an `$` secret key. So the ignore protected nothing; it only hid the one copy that is supposed to be canonical.

`scripts/deploy_testnet.sh` writes these files, so a redeploy now produces a diff. That is deliberate: a deployment that changes a contract id should leave a commit saying so, rather than silently editing an untracked file.

Both maps are committed as they stand today. I verified all five contracts across both networks against the backend's `.env.example` before committing — ten out of ten match, so this publishes the current truth rather than freezing a drift.

Unblocks the drift check in Orizon-Agents-BE-Stellar.

ALGOREX-PH added 3 commits 3 days ago

- stopped ignoring the deployed address book 6fac663
- tracked the testnet address book 2dc0692
- tracked the mainnet address book 036ec48

This was referenced 3 days ago

- 3.07 — Catch contract address drift** BIOcksmiths/Orizon-Agents-FE-Stellar#58 Closed
- 3.07 — Catch contract address drift** BIOcksmiths/Orizon-Agents-BE-Stellar#55 Closed
- Week 2 — Epic 3 backend (consolidates #50, #53, #54, #55)** BIOcksmiths/Orizon-Agents-BE-Stellar#56 Merged
- Week 2 — Epic 3 frontend (consolidates #53, #55, #57, #58)** BIOcksmiths/Orizon-Agents-FE-Stellar#59 Merged

ALGOREX-PH merged commit 88aa554 into `main` 3 days ago

Sign up for free to join this conversation on GitHub. Already have an account? [Sign in to comment](#)

The public reference external agent

`Orizon-Agents-Example-Agent-Stellar`, created 2026-09-15: a copyable external agent in one file (`agent.py`, one dependency) that **verifies the signature on Orizon's signed dispatches** before doing any work and answers with the response contract. Its README walks an operator through five steps — run, expose, bind, route, verify.

Source	https://github.com/BI0cksmiths/Orizon-Agents-Example-Agent-Stellar
Captured	2026-09-19 15:47 PHT · screenshots/12-reference-agent-repo.png · Shown: the top of the repository page; the rest of the README is in the full-page PNG.

Platform

Solutions

Resources

Open Source

Enterprise

Pricing

Search

Sign in

Sign up

BI0cksmiths / Orizon-Agents-Example-Agent-Stellar

Public

Notifications

Fork 0

Star 0

<> Code

Issues

Pull requests

Actions

Projects

Security and quality

Insights

main

4 Branches

0 Tags

Go to file

<> Code

ALGOREX-PH Merge pull request #1 from BI0cksmiths/fix/readme-and-pinni... 38a9510 · 3 days ago 53 Commits

.env.example	aligned the sample environment with the agent's rea...	3 days ago
.gitignore	added a python gitignore that keeps .env out and .en...	3 days ago
.python-version	pinned the python version render builds against	3 days ago
LICENSE	Initial commit	4 days ago
README.md	named the externally-managed-environment failure i...	3 days ago
agent.py	installed from the requirements file in the module do...	3 days ago
render.yaml	installed from the requirements file in the render blu...	3 days ago
requirements.txt	pinned pynacl to a supported range	3 days ago
test_agent.py	pinned pytest to a supported range in the test doctr...	3 days ago

README

MIT license

Orizon reference agent

A working external agent for the [Orizon](#) marketplace, in one file. Clone it, run it, put it on a public HTTPS URL, bind that URL to an agent you own on-chain, and the orchestrator will POST you paid work.

`agent.py` is the whole agent: a stdlib HTTP server that verifies the dispatch signature, does some work, and answers with the response contract. `pynacl` is its only dependency, pinned to a range in [requirements.txt](#); the Python it is built against is pinned in [.python-version](#), so a platform changing its default cannot break your build. There is no framework, no Docker, and no Orizon SDK.

Five commands. Do them in order.

#	Step	What it gets you
1	run	the agent answering on localhost
2	expose	a public HTTPS URL the orchestrator can reach
3	bind	that URL attached to an agent id you own
4	route	the planner putting a step on your agent
5	verify	proof the binding and the registration are real

Two things will cost you money or reputation if you skip them. Read [The one mistake that destroys your reputation](#) before you deploy, and [Getting paid](#) before you plan around revenue.

The protocol itself — the envelope, the signature, the freshness and replay rules — is documented once, in the backend's operator guide, and not restated here:

[Verifying an Orizon dispatch →](#)

About

A copyable reference external agent for the Orizon Agents network — one file, verifies the dispatch signature, runs on Stellar testnet.

Readme

MIT license

Activity

Custom properties

0 stars

0 watching

0 forks

Report repository

Releases

No releases published

Contributors 1

ALGOREX-PH

Danielle Meer

Languages

Python 100%

set_scorer on the ReputationLedger — on-chain rating writes enabled for production

2026-09-19

A testnet admin transaction on **2026-09-19** (06:16:57 UTC, ledger 4755006, Successful), invoked by GA7AI5...6R5OQV : `set_scorer(GDB4...CDHP)` on the ReputationLedger (CDCS...22ZT).

It makes the production backend's key (GDB4N25...) the ledger's authorized scorer, so **every paid run's ratings are written on-chain** by production. The `ratings` check shipped in `/readiness` (BE #59, Evidence 05) verifies that match on every boot.

Source <https://stellar.expert/explorer/testnet/tx/216e1b5f6ade4d75ec671bcda27b462bfd373d041b1ba2150d76002ee8d201f8>
Captured 2026-09-19 15:47 PHT · screenshots/13-set-scorer-tx-stellar-expert.png

The screenshot displays the StellarExpert interface for a transaction on the testnet. The transaction ID is 216e1b5f6ade4d75ec671bcda27b462bfd373d041b1ba2150d76002ee8d201f8. The transaction is successful and processed on 2026-09-19 at 06:16:57 UTC. It was invoked by account GA7AI5...6R5OQV, which called the `set_scorer` function on the ReputationLedger, setting the scorer to GDB4...CDHP. The transaction size is 516 bytes. The fee charged is 1.2143213 XLM, consisting of a non-refundable fee of 0.0027756 XLM and a refundable fee of 1.2115357 XLM, plus a 0.00001 XLM minimum operation fee. The transaction is valid before 2026-09-19 06:17:18. The interface also shows the transaction's preconditions and a list of signatures.

stellar expert NETWORK STATS ASSETS MARKETS SERVICES BLOG Network testnet

Transaction 216e1b5f6ade4d75ec671bcda27b462bfd373d041b1ba2150d76002ee8d201f8

Summary

Status: Successful
Ledger: 4755006
Source Account: GA7AI5...6R5OQV
Sequence Number: 8965185884651622
Transaction size: 516 bytes
Preconditions
Valid before: 2026-09-19 06:17:18

Processed: 2026-09-19 06:16:57 UTC
Max Fee: 1.3970109 XLM
Fee Charged: 1.2143213 XLM (0.0027756 non-refundable + 1.2115357 refundable + 0.00001 min op fee)

GA7AI5...6R5OQV invoked contract CDCS...22ZT `set_scorer(GDB4...CDHP)`

Signatures

GA7AI5...6R5OQV: LKRsGqJEbL7+4qc1TgV989G4B9I8GGHkRdQch0mSULL3vCTrvfuV3mMQhutu3d2qz2OYSfPPvk6XrpF/iIBQ==

2026 © StellarExpert v1.0.1
Open API docs Report a bug Dark theme
Terms of use Privacy policy Status
Binaries: CDQT5AG5V9GdZDPA7AIVAD0NIN346XUJOPZQNM27XNCCSTOQEIP

Week-2 public build post on X

Published 2026-09-19, 08:18 PHT

The Week-2 build thread from **@OrizonAgents402**, published **2026-09-19 at 08:18 PHT** (00:18 UTC), with a 40-second video: "This week, we focused on making external agents actually work safely and making reputation part of the routing".

x.com shows a blank page to a logged-out automated browser, so this capture is X's own embed view of the same post id (the second URL below).

Source	https://x.com/OrizonAgents402/status/2101103657043255772
Rendered via	https://platform.twitter.com/embed/Tweet.html?id=2101103657043255772
Captured	2026-09-19 15:47 PHT · screenshots/14-week-2-x-post.png

