

Orizon Agents — Week 2 Tranche Submission

Proof of Deliverables

Programme	Stellar Community Fund — Blue Belt Instawards (Cohort 2026)
Milestone	M2 · Week 2 — Reputation-Gated Routing (Deliverable D2) + External Agent Execution Path (Epic 2)
Sprint week	Mon 2026-09-14 → Fri 2026-09-18
Network	Stellar testnet only
Team	Danielle Bagaforo Meer — lead engineer (GitHub ALGOREX-PH) Rieselle Saure ("Rie") — PM + QA (GitHub rie-hash14)
Evidence captured	2026-09-19, from the live testnet deployment and public GitHub / Stellar Expert pages. Each page states its source URL and capture time.

DEPLOYED TESTNET CONTRACT IDS

AgentRegistry	CAPHXWU53UZUZJGV7IAE57NNMH3YYB5MTW06YA53KKMXSFVL0ITBJ3GQ
ReputationLedger	CDCS0BEVZUPQZV5GV4D6KYHZCLNGW2KXY74RUHSZ3EZUXF34DPW422ZT
PaymentEscrow	CBJPTMAPMG0DGZCZ2IMEQSRUX3WGUXNMKDTNN2KMJ3NFGYZ50J5525PI
AttestationRegistry	CBYUZK0ET43UXTBXZUJIBBJW50DGD2J2AZVVXCR3Q0NG0CAH0XQQHEGK
Asset SAC	CDLZFC3SYJYDZT7K67VZ75HPJVIEUVNIXF47ZG2FB2RMQQVU2HHG6CYSC

EVIDENCE IN THIS DOCUMENT

01	Plan card: each step's reputation and the routing floor, before payment	3.04 · D2	p. 2
02	Marketplace standing: floor stated once, provenance and standing marks	3.05 · D2	p. 3
03	Operator dashboard ("My Agents"), no wallet connected	2.06 · Epic 2	p. 4
04	Bind an endpoint to an on-chain agent, by wallet signature	2.01 / 2.05 · Epic 2	p. 5
05	Backend /readiness: cold-start headroom and whether ratings can be written	BE #58 / #59 · D2	p. 6
06	Backend PR #59 — planner-outage fallback, observable ratings, log redaction	BLO-121 · backend	p. 7
07	Frontend PR #62 — the plan card says when a plan is the planner fallback	BLO-121 · frontend	p. 8
08	UAT PR #3 — Rie's Week-2 QA commits	QA · rie-hash14	p. 9
09	Backend PRs merged 13–18 Sep 2026 (13)	backend delivery	p. 10
10	Frontend PRs merged 13–18 Sep 2026 (10)	frontend delivery	p. 11
11	Contracts PR #2 — the deployed address book is tracked in git	3.07	p. 12
12	The public reference external agent	2.04 · Epic 2	p. 13
13	set_scorer on the ReputationLedger — on-chain rating writes enabled for production	on-chain ratings enabled	p. 14
14	Week-2 public build post on X	7.02 · published 2026-09-19	p. 15

Every page names the public URL it was captured from, so each claim can be checked live. On the plan-card and marketplace pages every agent reads the ≈3.50 prior estimate: no on-chain ratings existed during Week 2, so those frames show the floor being stated and applied, not any agent being excluded by it.

Plan card: each step's reputation and the routing floor, before payment

The live orchestrator decomposed the demo intent **"tetris game in html"** (chosen with the page's own preset button, then *Decompose*) into a six-step plan. Before the buyer is asked to authorize anything, the card states the routing floor that was applied (**floor 2.75 · applied**, checked against each agent's reputation lower bound) and shows every step's agent with its reputation and that reputation's source: the **≈★ 3.50** chip, where **≈** marks a prior estimate rather than settled on-chain ratings.

Agents on the live registry carry the **≈3.50 starting estimate**, which clears the 2.75 floor, so the floor summary reads "the floor acted on no agents". The five agents it lists as having no endpoint bound are not candidates because there is nothing to dispatch a step to. Nothing was authorized, simulated or paid for this capture.

Source <https://orizons.xyz/app/orchestrator>

Captured 2026-09-19 14:59 PHT · screenshots/01-orizons-plan-card.png

The screenshot displays the Orizon Agents Orchestrator interface. The left sidebar contains a navigation menu with options: Overview, Agents, Register, Bind, My Agents, Reputation, Orchestrator (selected), Trace, Events, Send XLM, PDAX Ramp, Wallet, and Flow. The main area is titled 'Orchestrator' and shows the intent 'tetris game in html'. Below the intent, there are buttons for 'TETRIS GAME IN HTML', 'CALCULATOR WEB APP', 'SNAKE GAME IN HTML', and 'POMODORO TIMER WITH SOUND', along with a 'DECOMPOSE' button. The 'Execution plan' section shows a plan with 6 steps, a routing floor of 2.75 (APPLIED), and a total estimate of 0.168 XLM and 6.3s. The steps are:

Step	Agent	Reputation	Action	Estimate
01	RESEARCH.PRO	≈★ 3.50	extract feature brief + edge cases for the build	0.024 · 0.6s
02	SEO.BRIEF	≈★ 3.50	produce brand identity: name, tagline, audience	0.009 · 0.5s
03	DESIGN.FIGMA	≈★ 3.50	lock design tokens: palette, typography, motion	0.038 · 0.4s
04	CODE.GEN	≈★ 3.50	implement single-file HTML using brief + tokens	0.054 · 2.6s
05	CODE.CRITIC	≈★ 3.50	polish pass: a11y, motion, persistence, edge cases	0.052 · 1.8s
06	DEPLOY.VO	≈★ 3.50	seal artifact + record on-chain proof	0.011 · 0.4s

The 'REPUTATION FLOOR' section shows 'NO CHANGES' and '5 with no endpoint bound'. The 'WALLET REQUIRED' section shows 'Connect Freightier (testnet) to pay with x402 on-chain, or run a simulated pass.' and buttons for 'CONNECT WALLET', 'PAY WITH FIAT', and 'SIMULATE'.

Marketplace standing: floor stated once, provenance and standing marks

The Agent Registry states the selection floor **once**, above the table (**floor 2.75**, checked against each agent's reputation lower bound, never the headline score). Agents registered on-chain against the public registry carry an **"external"** provenance mark; the seeded first-party catalog rows carry none. Standing marks sit beside the agent: **"not yet operational"** (registered, no endpoint bound) and **"delisted by operator"**.

A third mark, **"below floor · not eligible"**, flags any agent whose lower bound falls under the floor; every agent on the live registry currently clears it.

Source <https://orizons.xyz/app/agents>

Captured 2026-09-19 14:59 PHT · screenshots/02-orizons-agents-standing.png · Shown in excerpts (slices of the full-page capture, 1440 × 2057 px); nothing inside a slice is altered.

ORIZON AGENTS

CONSOLE / AGENTS

CONNECT WALLET

LIVE

WORKSPACE

- Overview
- Agents
- Register
- Bind
- My Agents
- Reputation
- Orchestrator
- Trace
- Events
- Send XLM
- PDAX Ramp
- Wallet
- Flow

Agent Registry

ERC-8004 profiles — identity, skills, price, reputation.

+ REGISTER AGENT

Selection floor FLOOR 2.75

The floor decides which agents the orchestrator will consider when it builds a plan. It is checked against each agent's reputation lower bound — what a score is worth once it has been discounted for how few rated jobs stand behind it — and never against the headline score in the reputation column. An agent can show a strong score and still sit below the floor.

search name or skill...

ALL ROUTABLE ONLINE IDLE OFFLINE

ID	AGENT	SKILLS	PRICE / REPUTATION	CALL	RUNS	STATUS	
agt_01h8	copywrite.v3	COPY SEO EN	0.012	3.50	18,420	ONLINE	VIEW
agt_02k2	design.figma	UI TOKENS FIGMA	0.018	3.50	7,321	ONLINE	VIEW

... seeded catalog rows agt_03d9 → agt_12r0 omitted here (no marks) — all rows are in the full-page PNG ...

w1_audit_a7x	week1 audit scratch (non-executing)	AUDIT SCRATCH	0.001	3.50	0	ONLINE	VIEW
sign_probe_bb5c121.05	sign probe	PROBE SIGN	0.021	3.50	0	ONLINE	VIEW
Testing_Agent	Testing Agent		0.030	3.50	0	ONLINE	VIEW
dan_w1_probe	Dan Week 1 Registration	DEMO	0.050	3.50	0	OFFLINE	VIEW
spike_97437	Spike Agent	SPIKE	0.010	3.50	0	ONLINE	VIEW
uat605_ext_op	UAT 6.05 haiku operator	HAIKU POETRY UAT605	0.010	3.50	0	ONLINE	VIEW
calculatorai	Calculator AI	ADDING SUBTRACTING	0.010	3.50	0	ONLINE	VIEW
algorex	algorex		0.010	3.50	0	ONLINE	VIEW


Operator dashboard (“My Agents”), no wallet connected

The operator dashboard as a visitor sees it with **no wallet connected**. Agent ownership is read from the chain, so the page shows only a prompt to connect; connecting reads public registry data and signs nothing.

Once the owner connects their wallet, the same page lists every agent that wallet owns on-chain, with its reputation and routing standing, its endpoint binding, and a *Settlement* panel of what it has been paid (read from the escrow's on-chain charge events).

Source <https://orizons.xyz/app/operator>

Captured 2026-09-19 14:59 PHT · screenshots/03-orizons-operator-dashboard.png



Bind an endpoint to an on-chain agent, by wallet signature

The operator endpoint-binding page. An operator enters the id of the agent they registered on-chain and the HTTPS endpoint that runs its work; the registry issues a one-time challenge naming the agent, the endpoint and a nonce, and the **wallet that owns the agent signs it** — nothing else can authorize the binding. Ownership stays on-chain; the endpoint is an off-chain service record that can be re-bound.

Captured with no wallet connected, so the button reads “connect the owner wallet to sign”. `weather_bot` and `https://agent.example.com/run` are the form's placeholders, not a real binding.

Source <https://orizons.xyz/app/bind>

Captured 2026-09-19 14:59 PHT · screenshots/04-orizons-bind.png

ORIZON AGENTS

CONSOLE / BIND

CONNECT WALLET

LIVE

WORKSPACE

- Overview
- Agents
- Register
- Bind**
- My Agents
- Reputation
- Orchestrator
- Trace
- Events
- Send XLM
- PDAX Ramp
- Wallet
- Flow

2,498 agents online
avg completion 94%

operator
testnet

Bind an Endpoint

Point one of your agents at the HTTPS endpoint that runs its work. The wallet that owns the agent signs the binding — nothing else can authorize it.

CONNECT WALLET

ENDPOINT BINDING

SIGNING WALLET · CONNECT A WALLET

AGENT ID

weather_bot

the id you registered on-chain · letters, digits and underscore

CURRENT BINDING

enter an agent id to see what it points at today

ENDPOINT URL

https://agent.example.com/run

where the agent receives work · https is assumed when you omit the scheme

TRUST BOUNDARY

Ownership of this agent is on-chain and permanent. The endpoint is not: it is an off-chain service record held by the registry, so you can move hosts by binding again, and nothing about your ownership changes.

BIND ENDPOINT

connect the owner wallet to sign

HOW BINDING WORKS

1. The registry issues a one-time challenge naming your agent, the endpoint and a nonce.
2. Your wallet signs that exact string. The signature proves the binding came from the account that owns the agent — no password, no account to create.
3. The challenge is short-lived. If it expires while the wallet popup is open, a fresh one is requested rather than a dead signature submitted.

No agent yet? [Register one first](#)

Backend /readiness: cold-start headroom and whether ratings can be written

Captured 2026-09-19

The production backend's readiness probe. `cold_start` shows that a brand-new agent clears the routing floor: its reputation lower bound is **5677 bps against a 5500 bps floor — a 177 bps margin** (2.84 vs 2.75 on the dApp's 5-point scale), so `routable: true`. `ratings.writer` reports whether production can write ratings; it reads `scorer`: the backend's signer is the ReputationLedger's authorized scorer (`GDB4N25...`), so every paid run's ratings are written on-chain.

The JSON body was re-indented in the browser for legibility; its content is unchanged.

Source <https://orizon-agents-be-stellar.onrender.com/readiness>

Captured 2026-09-19 14:59 PHT · screenshots/05-be-readiness.png · Shown: the top 520 px of the 1440 × 900 frame; the rest of it is blank.

GET https://orizon-agents-be-stellar.onrender.com/readiness — fetched 2026-09-19T06:59:37.231Z (response body, indented for legibility)

```
{
  "status": "ready",
  "llm": "ok",
  "stellar": "configured",
  "signer": "configured",
  "pdax": "configured",
  "cold_start": {
    "routable": true,
    "lower_bound_bps": 5677,
    "floor_bps": 5500,
    "margin_bps": 177
  },
  "ratings": {
    "writer": "scorer",
    "signer": "GDB4N25UYM3YNTTAWX7LSGI2P70R62QZXRQWAGF5TFVENDKCTTCDHP",
    "scorer": "GDB4N25UYM3YNTTAWX7LSGI2P70R62QZXRQWAGF5TFVENDKCTTCDHP"
  }
}
```

Backend PR #59 — planner-outage fallback, observable ratings, log redaction

The backend half of the Week-2 branch `feat/week-2-dan` (124 commits, merged 2026-09-18). **BLO-121**: when the planner (LLM) fails, the buyer now gets a fallback plan built from agents that passed the routing checks, flagged `planner_fallback: true`, instead of a 502. **Observable ratings**: `/readiness` gains the `ratings` field, the service logs at boot whether it can write ratings, and a failed rating names its cause. **Log redaction**: configured secrets and key-shaped strings are masked on every log line.

Source <https://github.com/BIOcksmiths/Orizon-Agents-BE-Stellar/pull/59>
Captured 2026-09-19 14:59 PHT · screenshots/06-be-pr-59.png · Shown in excerpts (slices of the full-page capture, 1440 × 4818 px); nothing inside a slice is altered.

The screenshot shows a GitHub pull request page for the repository `BIOcksmiths/Orizon-Agents-BE-Stellar`. The pull request is titled "Week 2: planner outage fallback (BLO-121), observable ratings, log redaction #59" and is merged. It was created by `ALGOREX-PH` 19 hours ago, merging 124 commits from the `feat/week-2-dan` branch into the `main` branch. The page shows a conversation with one comment from `ALGOREX-PH` 19 hours ago, edited. The comment describes the changes made for BLO-121, including fixes for the planner outage, making ratings observable, and implementing log redaction. The comment is structured with a title "BLO-121: a planner outage now gets a modest plan, not a 502" followed by a list of bullet points detailing the bug, checks, timeouts, new response fields, provider text, and ratings observability. It also includes sections for "Ratings are observable" and "Secrets never reach a log line". The right sidebar shows no reviews, assignees, labels, projects, milestones, or development progress. The bottom of the page shows the merged commit `5822729` with one check passed.

Platform Solutions Resources Open Source Enterprise Pricing Search Sign in Sign up

BIOcksmiths / Orizon-Agents-BE-Stellar Public Notifications Fork 0 Star 1

<> Code Issues Pull requests 9 Actions Projects Security and quality Insights

Week 2: planner outage fallback (BLO-121), observable ratings, log redaction #59

Merged ALGOREX-PH merged 124 commits into main from feat/week-2-dan 19 hours ago

Conversation 0 Commits 124 Checks 1 Files changed 27 +2,875 -47

ALGOREX-PH commented 19 hours ago • edited Collaborator

Week 2 work for Dan: fixes **BLO-121**, and makes ratings observable, so a deployment that can't write ratings says so instead of failing silently.

BLO-121: a planner outage now gets a modest plan, not a 502

- The bug**: when the LLM call failed, agno put the provider's error `string` into `result.content`. `plan: Plan = result.content` then crashed, and the router returned `502 decompose_failed`.
- Checking the planner's answer**. `planner_plan()` only accepts a real `Plan` from a run agno didn't mark errored or cancelled.
 - An error string, a cancelled run, an unparsed answer, or a raised provider exception all degrade to the existing fallback plan.
 - That fallback comes from the agents offered to the planner, with `degraded` set by the floor rule.
- Timeouts are unchanged**: they still return `504 decompose_timeout`. Nothing offered is still `503 no_routable_agents`.
- New `DecomposeResponse.planner_fallback`** (additive, default false): true whenever the served plan is the fallback rather than the planner's own. The FE PR shows it to the buyer.
- The provider's text is never returned**. It is logged at `WARNING` with secrets redacted.
- Verified end to end with the ticket's own repro (blank key, "write a haiku about the sea"): `200`, `planner_fallback: true`.

Ratings are observable

The live ledger holds zero ratings, and nothing in the service could say why. The diagnosis is on BLO-41. Now:

- `/readiness` gains `ratings`**: `{"writer": "disabled|no_signer|scorer|not_scorer|unchecked", "signer": "...", "scorer": ...}`.
 - `scorer` is read from the ledger's instance storage and cached for 300 s, with failed reads retried after 30 s.
 - The probe itself makes no network call.
 - It is informational and never changes readiness.
 - Checked read-only against testnet: the local signer **is** the stored scorer.
- Startup line**: once at boot, the service logs whether it can write ratings. Anything other than `scorer` is logged at `WARNING`.
- No more silent skip**. `submit_ratings` used to `return` without a word when unconfigured. It now logs a rate-limited `WARNING` and traces `ratings not submitted: <reason>` on the task.
- Failed ratings name their cause**. A failed rating now reads `reputation submit failed for X: Unauthorized | NotFound | Replay | OutOfRange | contract error #N | rpc error`, and never echoes exception text.
- Two previously misreported outcomes fixed**: a rating whose transaction came back `FAILED` or timed out used to be traced as "rated N/100". It is now an error.

Secrets never reach a log line

- Masking**: `SecretRedactionLogFilter` on the root handler masks every configured secret by exact value, longest first. It also masks anything shaped like an API key (including provider-starred keys) or a Stellar secret seed, in messages, args and tracebacks. Public keys, contract ids and transaction hashes are untouched.
- agno's logs**: agno installs its own Rich handler with propagation off, so its lines bypassed the JSON format, the request id and any redaction. That includes the verbatim provider error it logs at `ERROR`. Its loggers are now routed through the root handler, at `WARNING`.
- One rule for the planner**: `loggable` now uses the same shared redaction.

... test-hygiene and verification notes and the 124-commit timeline omitted — all in the full-page PNG ...

ALGOREX-PH merged commit 5822729 into main 19 hours ago 1 check passed View details

Frontend PR #62 — the plan card says when a plan is the planner fallback

The frontend half of BLO-121 (34 commits, merged 2026-09-18). When the backend serves a fallback plan, the plan card shows a calm notice above Authorize and Simulate saying the planner was unavailable and the plan was built only from agents that passed the routing checks, with **"Ask the planner again"**, which re-runs the intent the buyer actually submitted. The description lists its own verification (unit, Playwright and axe checks); the merge commit shows 5 checks passed.

Source <https://github.com/BIOcksmiths/Orizon-Agents-FE-Stellar/pull/62>
Captured 2026-09-19 14:59 PHT · screenshots/07-fe-pr-62.png · Shown in excerpts (slices of the full-page capture, 1440 × 3317 px); nothing inside a slice is altered.

The screenshot shows a GitHub pull request interface for the repository 'BIOcksmiths/Orizon-Agents-FE-Stellar'. The pull request is titled 'Week 2: tell the buyer when a plan is the planner fallback (BLO-121) #62' and is in a 'Merged' state. It was merged by 'ALGOREX-PH' 19 hours ago, merging 34 commits from the 'feat/week-2-dan' branch into the 'main' branch. The interface includes navigation tabs for Code, Issues, Pull requests (8), Actions, Projects, Security and quality, and Insights. A summary bar shows 1 conversation, 34 commits, 4 checks, and 10 files changed, with a net change of +1,100 -18 lines. The main content area features a comment from 'ALGOREX-PH' (a collaborator) from 19 hours ago. The comment describes the frontend changes for BLO-121, explaining that a 'fallback plan' is shown when the planner is unavailable. It details the 'What changed' section, which includes updates to 'Types and guard', 'PlannerFallbackNotice' on the plan card, and the 'Ask the planner again' functionality. It also lists 'Tests' (Unit, E2E, and Playwright) and a 'Verification' section for the merged branch, showing successful checks for lint, vitest coverage (981/981), Playwright (105 passed, 0 failed), and build/bundle. The right sidebar shows sections for Reviewers, Assignees, Labels, Projects, Milestone, Development, and 1 participant. At the bottom, a merge commit summary shows 'ALGOREX-PH merged commit 3070f5c into main 19 hours ago' with 5 checks passed.

Platform Solutions Resources Open Source Enterprise Pricing Search Sign in Sign up

BIOcksmiths / Orizon-Agents-FE-Stellar Public Notifications Fork 0 Star 1

<> Code Issues Pull requests 8 Actions Projects Security and quality Insights

Week 2: tell the buyer when a plan is the planner fallback (BLO-121) #62

Merged ALGOREX-PH merged 34 commits into main from feat/week-2-dan 19 hours ago

Conversation 1 Commits 34 Checks 4 Files changed 10 +1,100 -18

ALGOREX-PH commented 19 hours ago Collaborator

The frontend half of BLO-121. When the backend serves a **fallback plan** because the planner failed or answered with nothing usable (`planner_fallback: true` , added in BE #59), the buyer is now told so before paying.

What changed

- **Types and guard.** `DecomposeResponse.planner_fallback?: boolean` . The guard accepts true, false, null or absent, and rejects anything else, in the same way it handles `reputation_degraded` .
- **PlannerFallbackNotice on the plan card** , above Authorize and Simulate:
 - A calm, informational notice (`role="status"` , never an alert). It uses violet, because magenta means a warning on this card, and a glyph plus text so the meaning never rests on colour alone.
 - It says the planner was unavailable, so this is a minimal fallback plan built only from agents that passed the routing checks.
 - It never names a cause and never discourages paying.
- **"Ask the planner again"** re-runs the page's existing decompose with the intent the buyer actually submitted, not whatever the intent box says by then. No new API. It is disabled while a simulate or authorize run is in progress.
- **Accessibility.** Authorize's `aria-describedby` lists every notice on the page, in reading order: this one, then the unverified-reputation banner. When neither is shown, the attribute is left off entirely.

Tests

- **Unit:** 29 new tests (notice, plan card, guard).
- **E2E:** new `e2e/plan-fallback.spec.ts` :
 - the notice appears above the actions, and doesn't appear when the flag is false or absent;
 - the describedby combinations;
 - axe WCAG 2.0/2.1 A+AA is clean in both states;
 - at 390px there's no sideways scroll and nothing is clipped;
 - the retry asks the same intent again, and the planner's plan replaces the fallback.

Verification (on the merged branch)

- `typecheck` , `lint` , and Prettier on tracked files: clean.
- `vitest --coverage` : **981/981**. Coverage 96.23 / 91.88 / 94.63 / 98.44 on `lib/**` .
- **Playwright: 105 passed, 0 failed** (from the summary line; no X).
- `build` , `bundle` (every route within 480 KB), `check:addresses` , `test:parity` 20/20, and `lighthouse` : pass.

... the 34-commit timeline and the Vercel preview comment omitted — both in the full-page PNG ...

ALGOREX-PH merged commit 3070f5c into main 19 hours ago 5 checks passed View details

UAT PR #3 — Rie’s Week-2 QA commits

Rie’s Week-2 QA work in the public UAT repository, merged as PR #3 on 2026-09-17: **142 commits, 141 of them by rie-hash14** (21 on 12 Sep, 50 on 16 Sep, 70 on 17 Sep) plus one merge of `main` by ALGOREX-PH. They add the UAT coverage for **6.02** (reputation floor), **6.05** (external dispatch) and **6.06** (operator surfaces), and the defect log entries **D-036 → D-049**.

Source	https://github.com/BI0cksmiths/Orizon-Agents-UAT-Stellar/pull/3/commits
Captured	2026-09-19 15:00 PHT · screenshots/08-uat-pr-3-rie-commits.png · Shown in excerpts (slices of the full-page capture, 1440 × 9837 px); nothing inside a slice is altered.

Platform · Solutions · Resources · Open Source · Enterprise · Pricing

Search / Sign in Sign up

BI0cksmiths / Orizon-Agents-UAT-Stellar Public

Notifications Fork 0 Star 0

<> Code Issues Pull requests Actions Projects Security and quality Insights

Uat #3

Merged ALGOREX-PH merged 142 commits into `main` from `uat` 2 days ago

Conversation 0 Commits 142 Checks 0 Files changed 25 +4,217 -6

Commits on Sep 12, 2026

updated uat test plan module
rie-hash14 committed last week
f75d289

updated uat test plan module
rie-hash14 committed last week
c75823e

updated uat test plan module
rie-hash14 committed last week
bda5a74

updated uat test plan module
rie-hash14 committed last week
57ca660

added uat lockfile module
rie-hash14 committed last week
b3df188

updated uat traceability module
rie-hash14 committed last week
f68bf37

... remaining 15 commits of 12 Sep omitted ...

Commits on Sep 16, 2026

updated reputation floor spec module
rie-hash14 committed 3 days ago
a461492

updated uat defects module
rie-hash14 committed 3 days ago
2849ce5

updated uat defects module
rie-hash14 committed 3 days ago
3bf60a7

... remaining 47 commits of 16 Sep omitted ...

Commits on Sep 17, 2026

updated tsconfig module
rie-hash14 committed 2 days ago
903d916

added stellar key module
rie-hash14 committed 2 days ago
35aeb8b

added verify dispatch module
rie-hash14 committed 2 days ago
3ef88de

... the 17 Sep commits in between omitted — every commit is in the full-page PNG ...

merged operator surfaces spec module
rie-hash14 committed 2 days ago
700abc9

Merge branch 'main' into uat
ALGOREX-PH authored 2 days ago
Verified a46523c

Orizon Agents — Week 2 Tranche Submission · Proof of Deliverables · Stellar testnet

9 / 15

Backend PRs merged 13–18 Sep 2026 (13)

GitHub's own search for backend PRs merged 13–18 Sep 2026 returns **13**, numbered #41 to #59, covering Epic 2 (2.01 endpoint binding, 2.02 verifiable bounded dispatch, 2.03 external step failure semantics, 2.04 reference-agent contract, 2.06 settlement evidence), Epic 3 (3.02 excluded agents and the floor, 3.03 routing-time reads) and the Week-2 branch.

Source	https://github.com/BI0cksmiths/Orizon-Agents-BE-Stellar/pulls?q=is%3Apr+is%3Amerged+merged%3A2026-09-13..2026-09-18
Captured	2026-09-19 15:00 PHT · screenshots/09-be-pull-requests.png

PlatformSolutionsResourcesOpen SourceEnterprisePricing

Search / Sign in Sign up

BI0cksmiths / Orizon-Agents-BE-Stellar Public

Notifications Fork 0 Star 1

<> Code Issues Pull requests 9 Actions Projects Security and quality Insights

Q is:pr is:merged merged:2026-09-13..2026-09-18

Labels 12 Milestones 0 New pull request

Clear current search query, filters, and sorts

13 Total

AuthorLabelProjectsMilestonesReviewsAssigneeSort

Week 2: planner outage fallback (BLO-121), observable ratings, log redaction

#59 by ALGOREX-PH Collaborator was merged 19 hours ago

Epic 3 hardening: hold the floor after planning, enrich plan steps, ops fixes

#58 by ALGOREX-PH Collaborator was merged yesterday

Epic 2 hardening — kill switch, dispatch security, rating semantics

#57 by ALGOREX-PH Collaborator was merged 3 days ago

Week 2 — Epic 3 backend (consolidates #50, #53, #54, #55)

#56 by ALGOREX-PH Collaborator was merged 3 days ago

3.03 — Routing-time read hardening and degradation visibility

#52 by ALGOREX-PH Collaborator was merged 3 days ago

3.02 — Surface excluded agents and the floor

#51 by ALGOREX-PH Collaborator was merged 3 days ago

2.06 — Settlement evidence endpoint

#49 by ALGOREX-PH Collaborator was merged 3 days ago

Story 2.04: reference agent contract, published deadline, and the settlement finding

#48 by ALGOREX-PH Collaborator was merged 3 days ago

Story 2.03: external step failure semantics and rating consequences

#47 by ALGOREX-PH Collaborator was merged 4 days ago

Story 2.02: verifiable, bounded external dispatch

#45 by ALGOREX-PH Collaborator was merged 4 days ago

Harden the external step path: crash, rating and planner-injection fixes

#44 by ALGOREX-PH Collaborator was merged 4 days ago

Story 2.01: bind an operator endpoint to an on-chain agent id

#43 by ALGOREX-PH Collaborator was merged 4 days ago

Backend production hardening: pin money-route auth + guard consistency

#41 by ALGOREX-PH Collaborator was merged 4 days ago

ProTip! Mix and match filters to narrow down what you're looking for.

© 2026 GitHub, Inc.

TermsPrivacySecurityStatusCommunityDocsContactManage cookiesDo not share my personal information

Orizon Agents — Week 2 Tranche Submission · Proof of Deliverables · Stellar testnet

10 / 15

Frontend PRs merged 13–18 Sep 2026 (10)

The same search on the frontend returns **10**, numbered #48 to #62, including the 2.01 binding UI (#50), the 2.05 binding step in registration (#51), the 2.06 operator dashboard (#52), the 3.02 floor visibility contract (#54), the 3.04 plan card (#56) and the Week-2 branch (#62).

Source	https://github.com/BIOcksmiths/Orizon-Agents-FE-Stellar/pulls?q=is%3Apr+is%3Amerged+merged%3A2026-09-13..2026-09-18
Captured	2026-09-19 15:00 PHT · screenshots/10-fe-pull-requests.png

Platform · Solutions · Resources · Open Source · Enterprise · Pricing

Search

Sign in

Sign up

BIOcksmiths/Orizon-Agents-FE-Stellar

Public

Notifications

Fork 0

Star 1

<> Code

Issues

Pull requests 8

Actions

Projects

Security and quality

Insights

Q is:pr is:merged merged:2026-09-13..2026-09-18

Labels 12

Milestones 0

New pull request

Clear current search query, filters, and sorts

10 Total

Author · Label · Projects · Milestones · Reviews · Assignee · Sort

Week 2: tell the buyer when a plan is the planner fallback (BLO-121)

#62 by ALGOREX-PH Collaborator was merged 19 hours ago

1

Epic 3 hardening: green CI, honest reputation surfaces, live contract parity

#61 by ALGOREX-PH Collaborator was merged yesterday

1

Epic 2 hardening — rebind coverage and transaction links

#60 by ALGOREX-PH Collaborator was merged 3 days ago

1

Week 2 — Epic 3 frontend (consolidates #53, #55, #57, #58)

#59 by ALGOREX-PH Collaborator was merged 3 days ago

1

3.04 — Plan card: reputation, source and exclusions

#56 by ALGOREX-PH Collaborator was merged 3 days ago

1

3.02 — Floor visibility contract

#54 by ALGOREX-PH Collaborator was merged 3 days ago

1

2.06 — Operator dashboard

#52 by ALGOREX-PH Collaborator was merged 3 days ago

1

Story 2.05: endpoint binding step in the registration flow

#51 by ALGOREX-PH Collaborator was merged 3 days ago

1

Story 2.01: operator endpoint binding UI

#50 by ALGOREX-PH Collaborator was merged 4 days ago

1

Frontend production hardening: a11y, skills-input UX, SSE guard

#48 by ALGOREX-PH Collaborator was merged 4 days ago

1

ProTip!

What's not been updated in a month: [updated:<2026-08-19](#).

© 2026 GitHub, Inc. Terms Privacy Security Status Community Docs Contact Manage cookies Do not share my personal information

Contracts PR #2 — the deployed address book is tracked in git

Merged 2026-09-16. The contract repo's record of what is deployed (`addresses.json` for testnet, `addresses.mainnet.json`) was git-ignored, so it existed only on the machine that last deployed. It is now **tracked in git**, giving story 3.07's contract-id drift check a canonical file to compare the backend configuration against. The PR records that all five contract ids on both networks matched the backend's `.env.example` when committed.

Source <https://github.com/BIOcksmiths/Orizon-Agents-Smart-Contract-Stellar/pull/2>
Captured 2026-09-19 15:00 PHT · screenshots/11-contracts-pr-2.png

The screenshot shows a GitHub pull request (PR) titled "Track the deployed address book #2" in the repository "BIOcksmiths/Orizon-Agents-Smart-Contract-Stellar". The PR is merged and was created 3 days ago. It shows a conversation with a comment from ALGOREX-PH explaining the need to track the deployed address book in git. The comment includes details about the files being tracked and the reason for the change. The PR history shows three commits: "stopped ignoring the deployed address book", "tracked the testnet address book", and "tracked the mainnet address book". The PR is referenced by two other PRs: "3.07 — Catch contract address drift" and "Week 2 — Epic 3 backend (consolidates #50, #53, #54, #55)".

Platform Solutions Resources Open Source Enterprise Pricing Search / Sign in Sign up

BIOcksmiths / Orizon-Agents-Smart-Contract-Stellar Public

Code Issues Pull requests Actions Projects Security and quality Insights

Track the deployed address book #2

Merged ALGOREX-PH merged 3 commits into main from feat/3.07-track-addresses 3 days ago

Conversation 0 Commits 3 Checks 0 Files changed 3 +38 -2

ALGOREX-PH commented 3 days ago Collaborator

Story 3.07 (BLO-111) asks for a check that catches contract-id drift between this repo and the backend config. That check could not exist, because **the file it would check against was not in this repo.**

```
38 tracked files · addresses tracked: 0
.gitignore:7 addresses.json
.gitignore:8 addresses.*.json
```

`addresses.json` and `addresses.mainnet.json` are this repo's record of which contracts are deployed, and ignoring them meant that record existed only on whichever laptop last ran a deploy. A fresh clone had no address book at all. No CI anywhere could compare against one. And the authoritative copy was one disk failure away from being recoverable only by reading it back out of a *downstream* copy in the backend's `.env.example`.

The ticket says the drift is "invisible to review and to CI by construction". This is the construction.

Nothing in these files is secret. Contract ids are public on-chain identifiers; the `admin` field is a public G-address. Every value is already committed to `README.md` here, and to `.env.example`, `README.md`, `onchain-details.tsx` and `lib/api.test.ts` in the other two repos — verified before committing, along with a check that neither file contains an `$` secret key. So the ignore protected nothing; it only hid the one copy that is supposed to be canonical.

`scripts/deploy_testnet.sh` writes these files, so a redeploy now produces a diff. That is deliberate: a deployment that changes a contract id should leave a commit saying so, rather than silently editing an untracked file.

Both maps are committed as they stand today. I verified all five contracts across both networks against the backend's `.env.example` before committing — ten out of ten match, so this publishes the current truth rather than freezing a drift.

Unblocks the drift check in Orizon-Agents-BE-Stellar.

ALGOREX-PH added 3 commits 3 days ago

- stopped ignoring the deployed address book 6fac663
- tracked the testnet address book 2dc0692
- tracked the mainnet address book 036ec48

This was referenced 3 days ago

- 3.07 — Catch contract address drift BIOcksmiths/Orizon-Agents-FE-Stellar#58 Closed
- 3.07 — Catch contract address drift BIOcksmiths/Orizon-Agents-BE-Stellar#55 Closed
- Week 2 — Epic 3 backend (consolidates #50, #53, #54, #55) BIOcksmiths/Orizon-Agents-BE-Stellar#56 Merged
- Week 2 — Epic 3 frontend (consolidates #53, #55, #57, #58) BIOcksmiths/Orizon-Agents-FE-Stellar#59 Merged

ALGOREX-PH merged commit 88aa554 into main 3 days ago

Sign up for free to join this conversation on GitHub. Already have an account? [Sign in to comment](#)

© 2026 GitHub, Inc. Terms Privacy Security Status Community Docs Contact Manage cookies Do not share my personal information

The public reference external agent

`Orizon-Agents-Example-Agent-Stellar`, created 2026-09-15: a copyable external agent in one file (`agent.py`, one dependency) that **verifies the signature on Orizon's signed dispatches** before doing any work and answers with the response contract. Its README walks an operator through five steps — run, expose, bind, route, verify.

Source	https://github.com/BI0cksmiths/Orizon-Agents-Example-Agent-Stellar
Captured	2026-09-19 15:00 PHT · screenshots/12-reference-agent-repo.png · Shown: the top of the repository page; the rest of the README is in the full-page PNG.

Platform

Solutions

Resources

Open Source

Enterprise

Pricing

Search

Sign in

Sign up

BI0cksmiths / Orizon-Agents-Example-Agent-Stellar

Public

Notifications

Fork 0

Star 0

<> Code

Issues

Pull requests

Actions

Projects

Security and quality

Insights

main

4 Branches

0 Tags

Go to file

<> Code

ALGOREX-PH Merge pull request #1 from BI0cksmiths/fix/readme-and-pinni... 38a9510 · 3 days ago 53 Commits

.env.example	aligned the sample environment with the agent's rea...	3 days ago
.gitignore	added a python gitignore that keeps .env out and .en...	3 days ago
.python-version	pinned the python version render builds against	3 days ago
LICENSE	Initial commit	4 days ago
README.md	named the externally-managed-environment failure i...	3 days ago
agent.py	installed from the requirements file in the module do...	3 days ago
render.yaml	installed from the requirements file in the render blu...	3 days ago
requirements.txt	pinned pynacl to a supported range	3 days ago
test_agent.py	pinned pytest to a supported range in the test doctr...	3 days ago

README

MIT license

Orizon reference agent

A working external agent for the [Orizon](#) marketplace, in one file. Clone it, run it, put it on a public HTTPS URL, bind that URL to an agent you own on-chain, and the orchestrator will POST you paid work.

`agent.py` is the whole agent: a stdlib HTTP server that verifies the dispatch signature, does some work, and answers with the response contract. `pynacl` is its only dependency, pinned to a range in [requirements.txt](#); the Python it is built against is pinned in [.python-version](#), so a platform changing its default cannot break your build. There is no framework, no Docker, and no Orizon SDK.

Five commands. Do them in order.

#	Step	What it gets you
1	run	the agent answering on localhost
2	expose	a public HTTPS URL the orchestrator can reach
3	bind	that URL attached to an agent id you own
4	route	the planner putting a step on your agent
5	verify	proof the binding and the registration are real

Two things will cost you money or reputation if you skip them. Read [The one mistake that destroys your reputation](#) before you deploy, and [Getting paid](#) before you plan around revenue.

The protocol itself — the envelope, the signature, the freshness and replay rules — is documented once, in the backend's operator guide, and not restated here:

[Verifying an Orizon dispatch →](#)

About

A copyable reference external agent for the Orizon Agents network — one file, verifies the dispatch signature, runs on Stellar testnet.

Readme

MIT license

Activity

Custom properties

0 stars

0 watching

0 forks

Report repository

Releases

No releases published

Contributors 1

ALGOREX-PH

Danielle Meer

Languages

Python 100%

set_scorer on the ReputationLedger — on-chain rating writes enabled for production

2026-09-19

A testnet admin transaction on **2026-09-19** (06:16:57 UTC, ledger 4755006, Successful), invoked by GA7AI5...6R5OQV : `set_scorer(GDB4...CDHP)` on the ReputationLedger (CDCS...22ZT).

It makes the production backend's key (GDB4N25...) the ledger's authorized scorer, so **every paid run's ratings are written on-chain** by production. The `ratings` check shipped in `/readiness` (BE #59, Evidence 05) verifies that match on every boot.

Source
Captured

<https://stellar.expert/explorer/testnet/tx/216e1b5f6ade4d75ec671bcda27b462bfd373d041b1ba2150d76002ee8d201f8>
2026-09-19 15:00 PHT · screenshots/13-set-scorer-tx-stellar-expert.png

The screenshot displays the StellarExpert interface for a transaction on the testnet. The transaction ID is 216e1b5f6ade4d75ec671bcda27b462bfd373d041b1ba2150d76002ee8d201f8. The summary section indicates a successful transaction on ledger 4755006, processed at 2026-09-19 06:16:57 UTC. The source account is GA7AI5...6R5OQV, and the fee charged is 1.2143213 XLM. The transaction invokes the `set_scorer` contract on the ReputationLedger, setting the scorer to GDB4...CDHP. The interface also shows the transaction size (516 bytes) and the valid before time (2026-09-19 06:17:18). The footer includes the StellarExpert version (v1.0.1) and links to API docs, bug reports, and terms of use.

Week-2 public build post on X

Published 2026-09-19, 08:18 PHT

The Week-2 build thread from **@OrizonAgents402**, published **2026-09-19 at 08:18 PHT** (00:18 UTC), with a 40-second video: "This week, we focused on making external agents actually work safely and making reputation part of the routing".

x.com shows a blank page to a logged-out automated browser, so this capture is X's own embed view of the same post id (the second URL below).

Source	https://x.com/OrizonAgents402/status/2101103657043255772
Rendered via	https://platform.twitter.com/embed/Tweet.html?id=2101103657043255772
Captured	2026-09-19 15:13 PHT · screenshots/14-week-2-x-post.png

